

GITHUB COPILOT

Training

Hands-on course

20
25.

Presented by
Maria Sukhareva



Website
www.mariasukhareva.com

Table of Contents

- C Foundations - 45 Min**
- C Automation Paradigm & Tooling - 60 Min**
- C Practical Lab (VS Code) - 4 h**
- C Governance and Scale-out - 30 Min**
- C Wrap-Up & Next Steps - 20 Min**

Foundations

Date: June 11, 2030



AI-Assisted Coding Motivation

Multiple studies have shown that AI-assisted coding increases both productivity and the quality of code



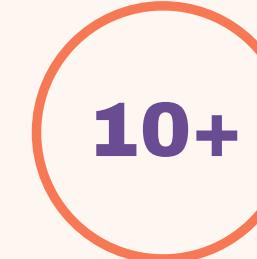
developers say AI had a positive effect on productivity¹



increased likelihood of passing all unit tests³



developers are using AI⁵



hours saved per week as reported by developers²

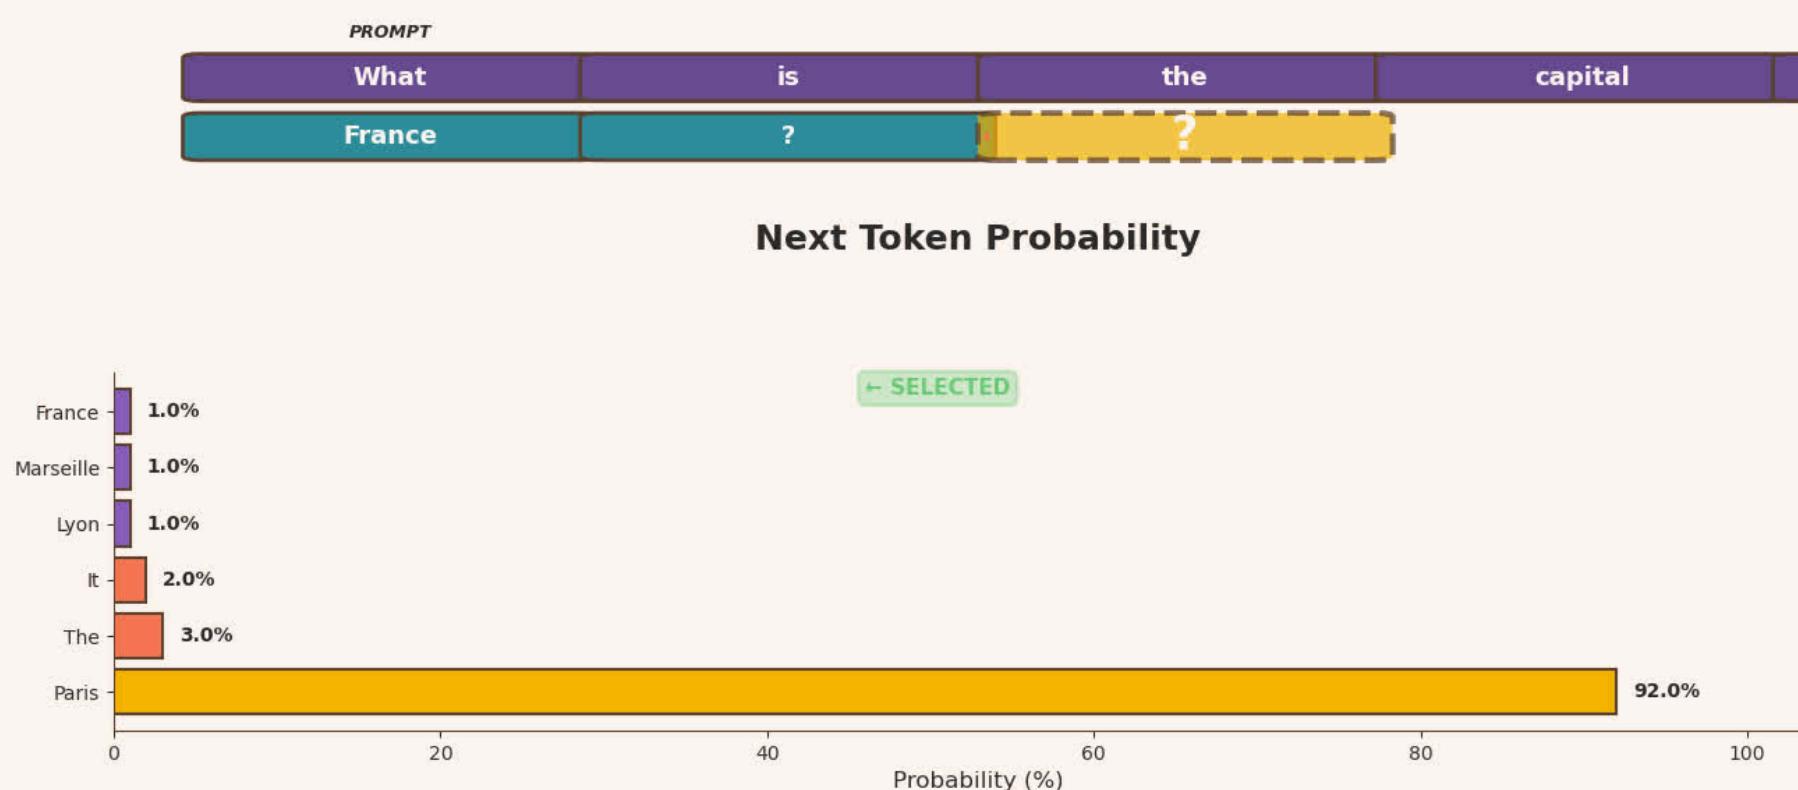


productivity boost⁴

References:

1. Stack Overflow Developer Survey 2025
2. Atlassian 2025 report
3. GitHub blog
4. IEEE Spectrum (AMD engineering op-ed, 2025)
5. DORA 2025

LLMs: Next Token Predictors



Large Language Models Have One Job

**Predict the next Token Based
On Provided Context**

- LLMs do not have a communicative intent (e.g., inform, lie).
- LLMs do not reason.
- They rely on patterns and token probabilities.

“ Critics call LLMs “stochastic parrots” because they sample next tokens from learned probabilities, not understanding. ”



Why LLMs Are Great For Programming Languages

Artificial vs. Natural Languages

Key differences between programming and natural languages

Precision



Designed to be unambiguous so machines parse them reliably

Fixed, Immutable Grammar



have stable, explicit grammars with (virtually) no exceptions.

Objective “Gold Standard”

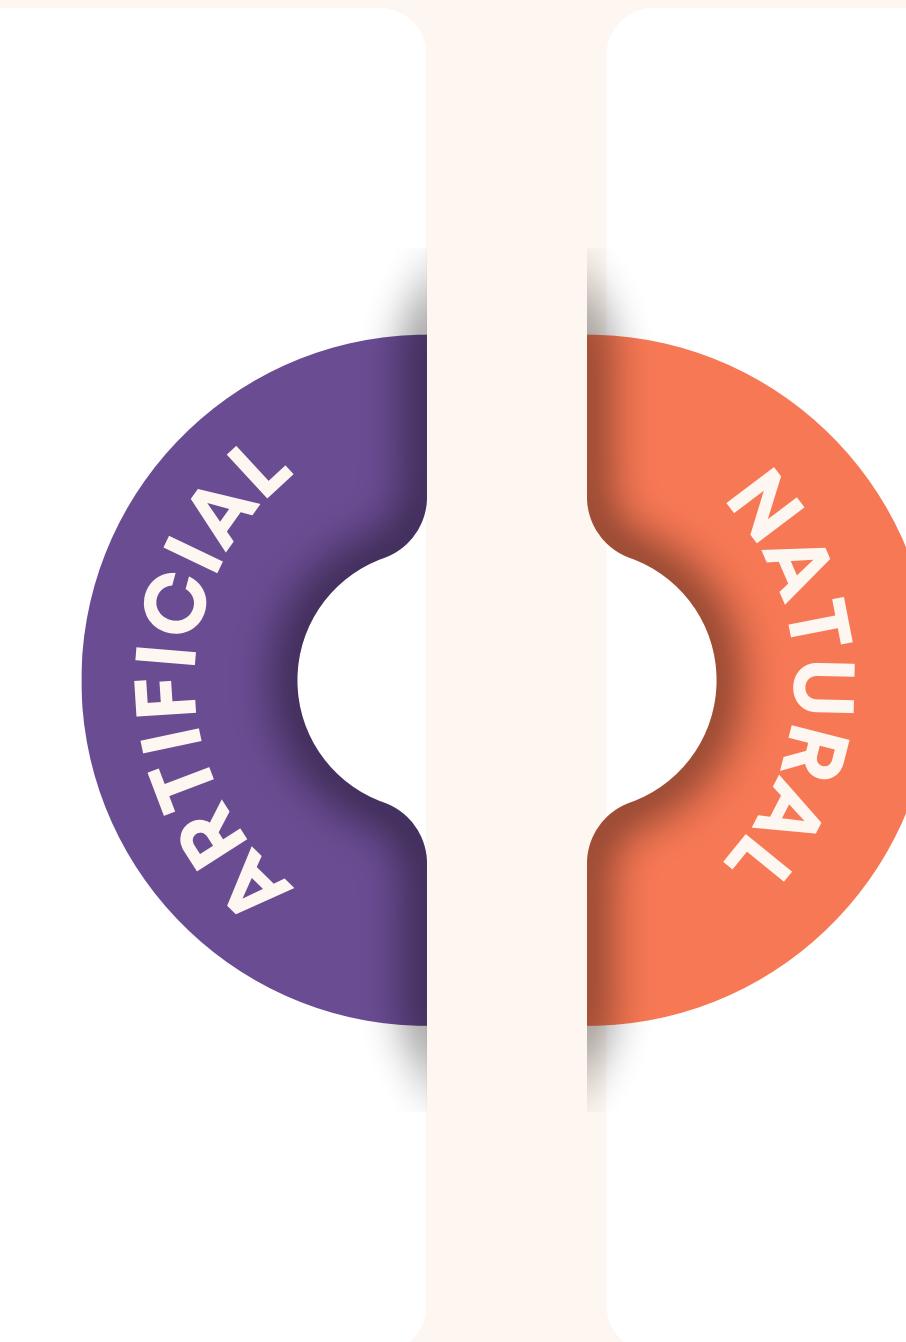


has objectively correct/optimal forms that can be evaluated and tested.

Creativity Penalized



Sampling randomness (“vibe coding”) quickly degrades correctness and maintainability



Ambiguity



Ambiguous on every level and rely on common sense

Evolving, Exception-ridden



evolve and team with irregularities and exceptions

Many Valid Phrasings



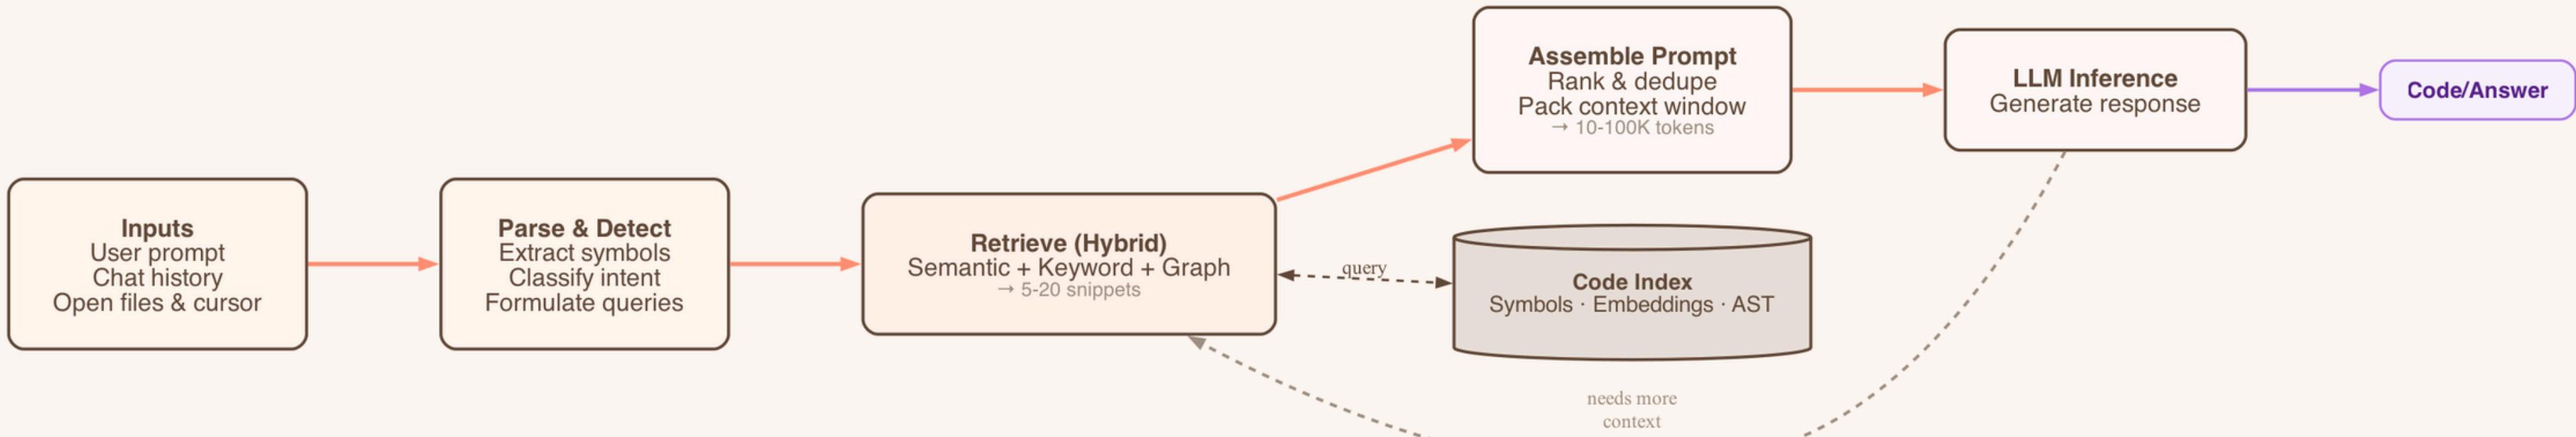
allows multiple equally acceptable ways to say the same thing

Creativity Tolerated



can be non-deterministic and varied without harm

How Coding Copilots Work



Retrieval Augmented Generation

The workspace is saved into a search index



Question One

Why is RAG needed here?



The Assembled Prompt Contains

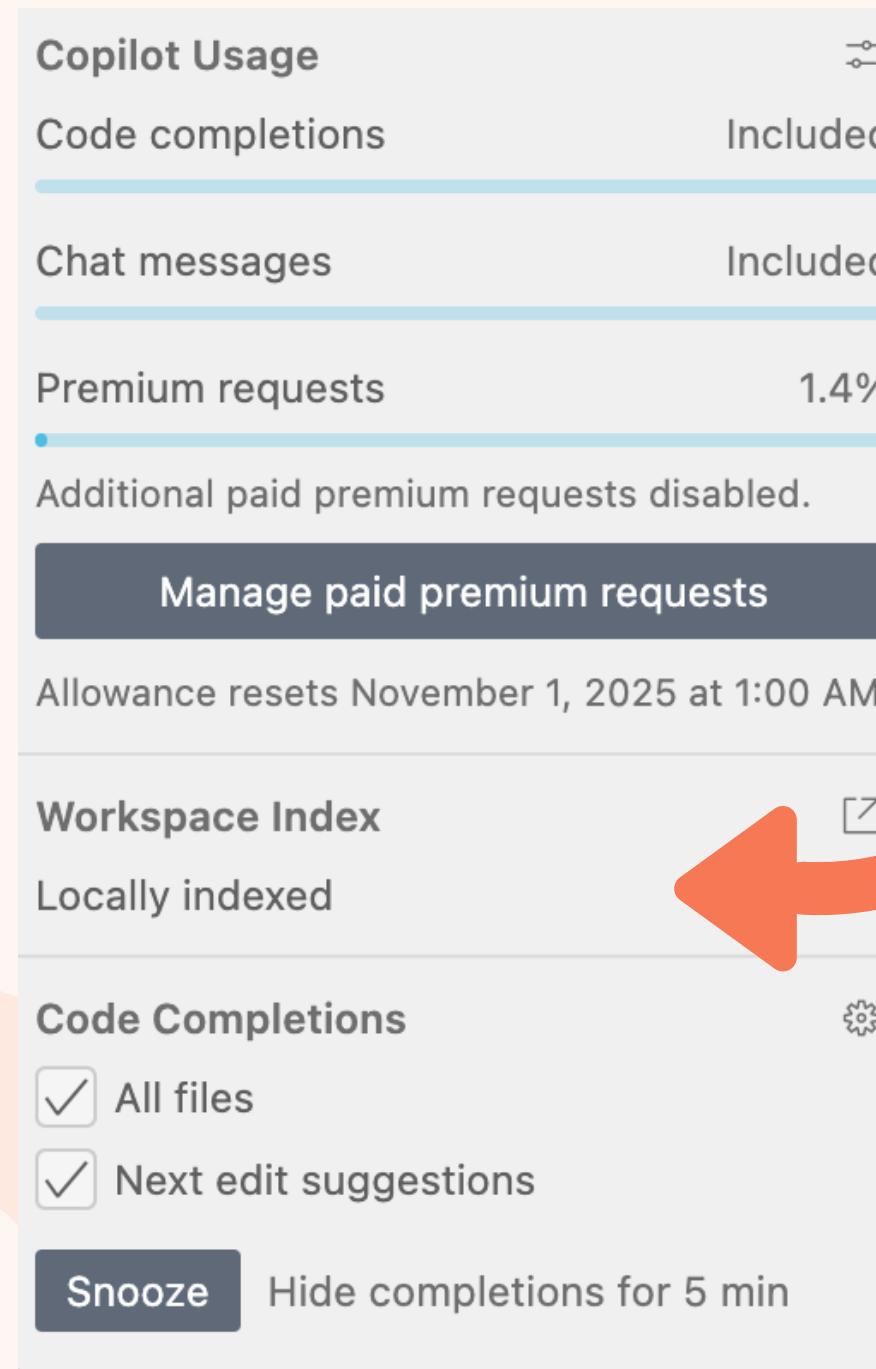
Code snippets, chat history, graph, repo description, user query etc.



Question Two

Why is long context window of the model essential for coding assistance?

Check Indexing Status



Check Status

Bottom bar → click Copilot item → Status/Details.

Create Local Workspace Index

- ⌘↑P (mac) / Ctrl+Shift+P (win/linux)
- Chat: Build Remote Workspace Index
- Check Status Again



Best Models In Github Copilot

Claude Sonnet 4.5

Cost: \$8.26, Multiplier: 1
SWE-bench: 69.8%

Reasoning heavy, deep refactor

01

02

03

04

05

06

07

08

Grok Code Fast

Cost: \$1.96, Multiplier: 0.25
SWE-bench: 52.4%
Fast coding, big context

GPT-5 Codex

Cost: \$2.21, Multiplier: 1
SWE-bench: 69.4%

Multi-file refactor and code review

GPT-5

Cost: \$1.41, Multiplier: 1
SWE-bench: 69.4%

Comparable with GPT-5-Codex

Gemini 2.5 Pro

Cost: \$0.38, Multiplier: 1
SWE-bench: 55.6%
Long context precision

GPT-5-mini

Cost: \$0.94, Multiplier: 0 (free)
SWE-bench: 59.2%
Fast coding, simple tasks

GPT-4.1 (default)

Cost: \$0.45, Multiplier: 0
SWE-bench: 47.4%
Fast coding, simple tasks

GPT-4o

Cost: \$1.53, Multiplier 0
SWE-bench: 27.2%
Very simple tasks



Claude Opus has multiplier 10!

Cost can be found here

Best Models In Github Copilot

Hey Coders, what do you like more in your coding copilot?

You can see how people vote. [Learn more](#)

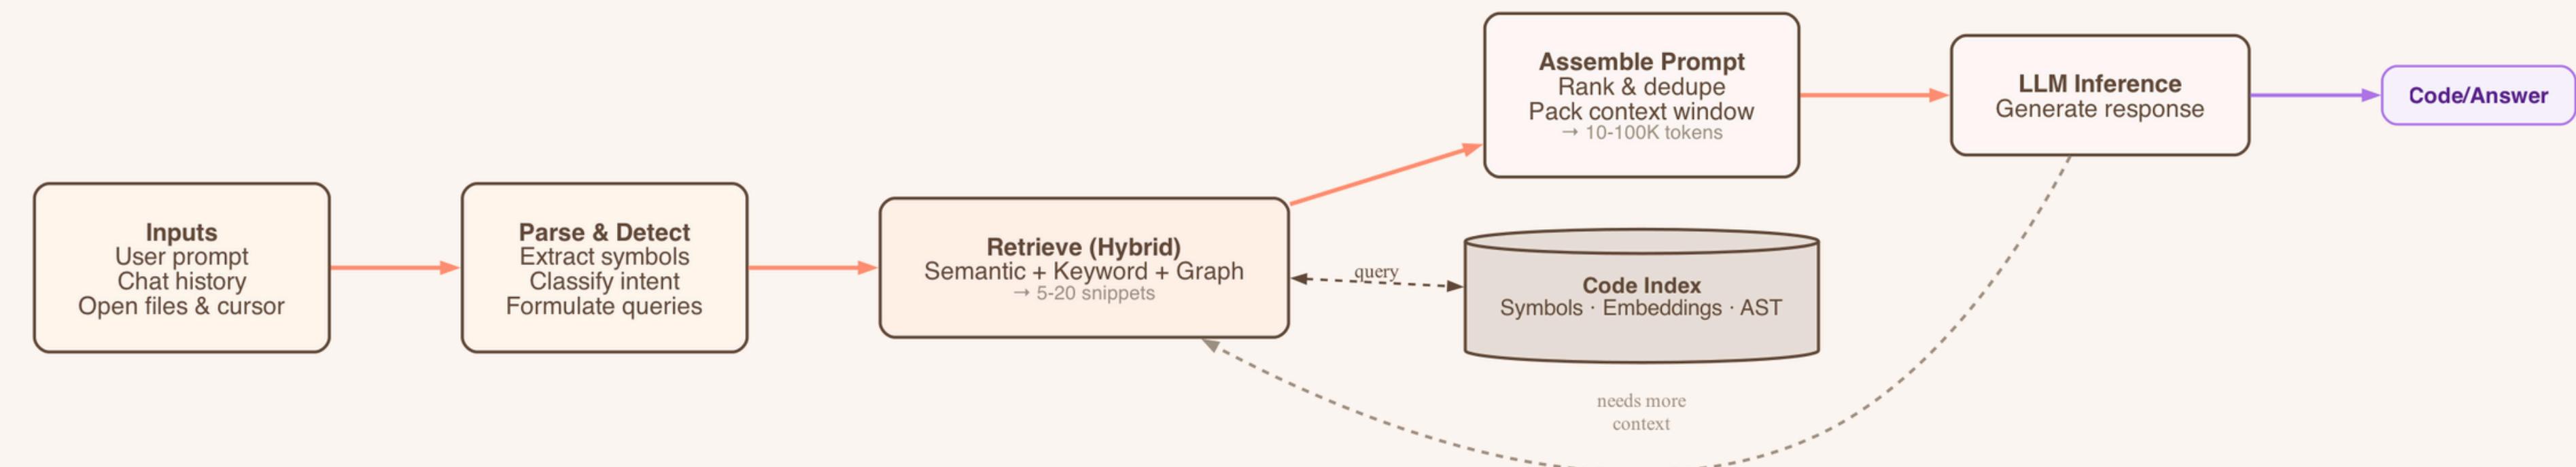


[349 votes](#) • 4d left • [Undo](#)

Just try and decided yourself!

Search vs. LLM

What is more important?



Search is The Key



Grounding

A short summary of the fundamental decisions and solution strategies



Task Focus

They are the basis for many other detailed decisions or implementation rules.



Token Economy

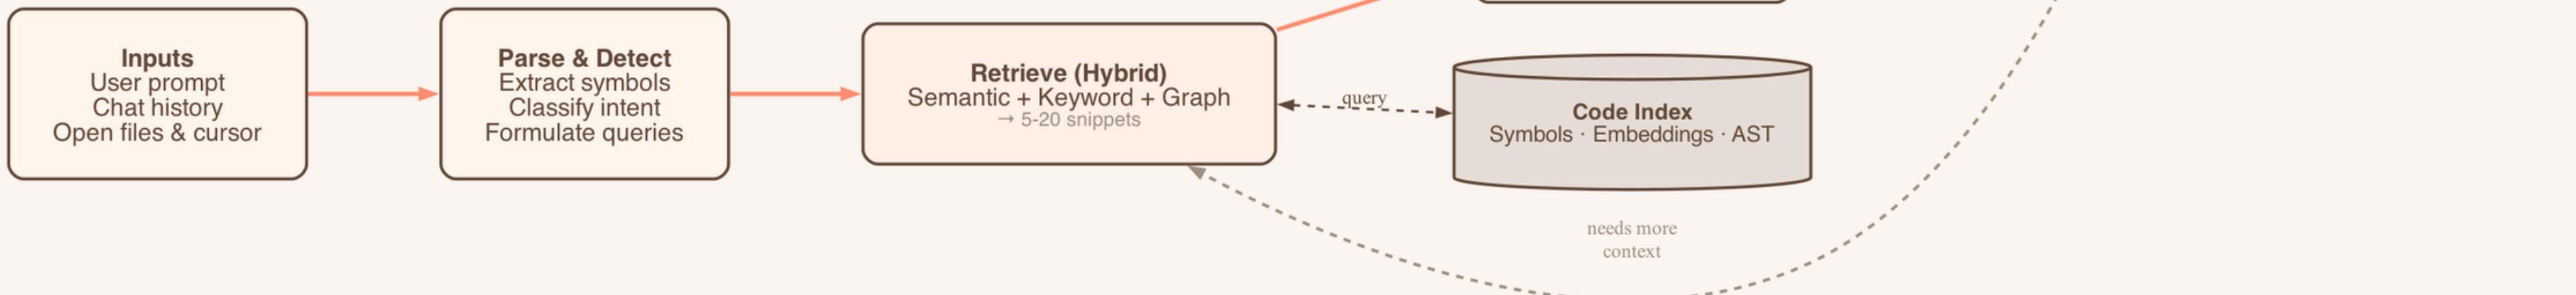
Keep the explanation of these key decisions very short



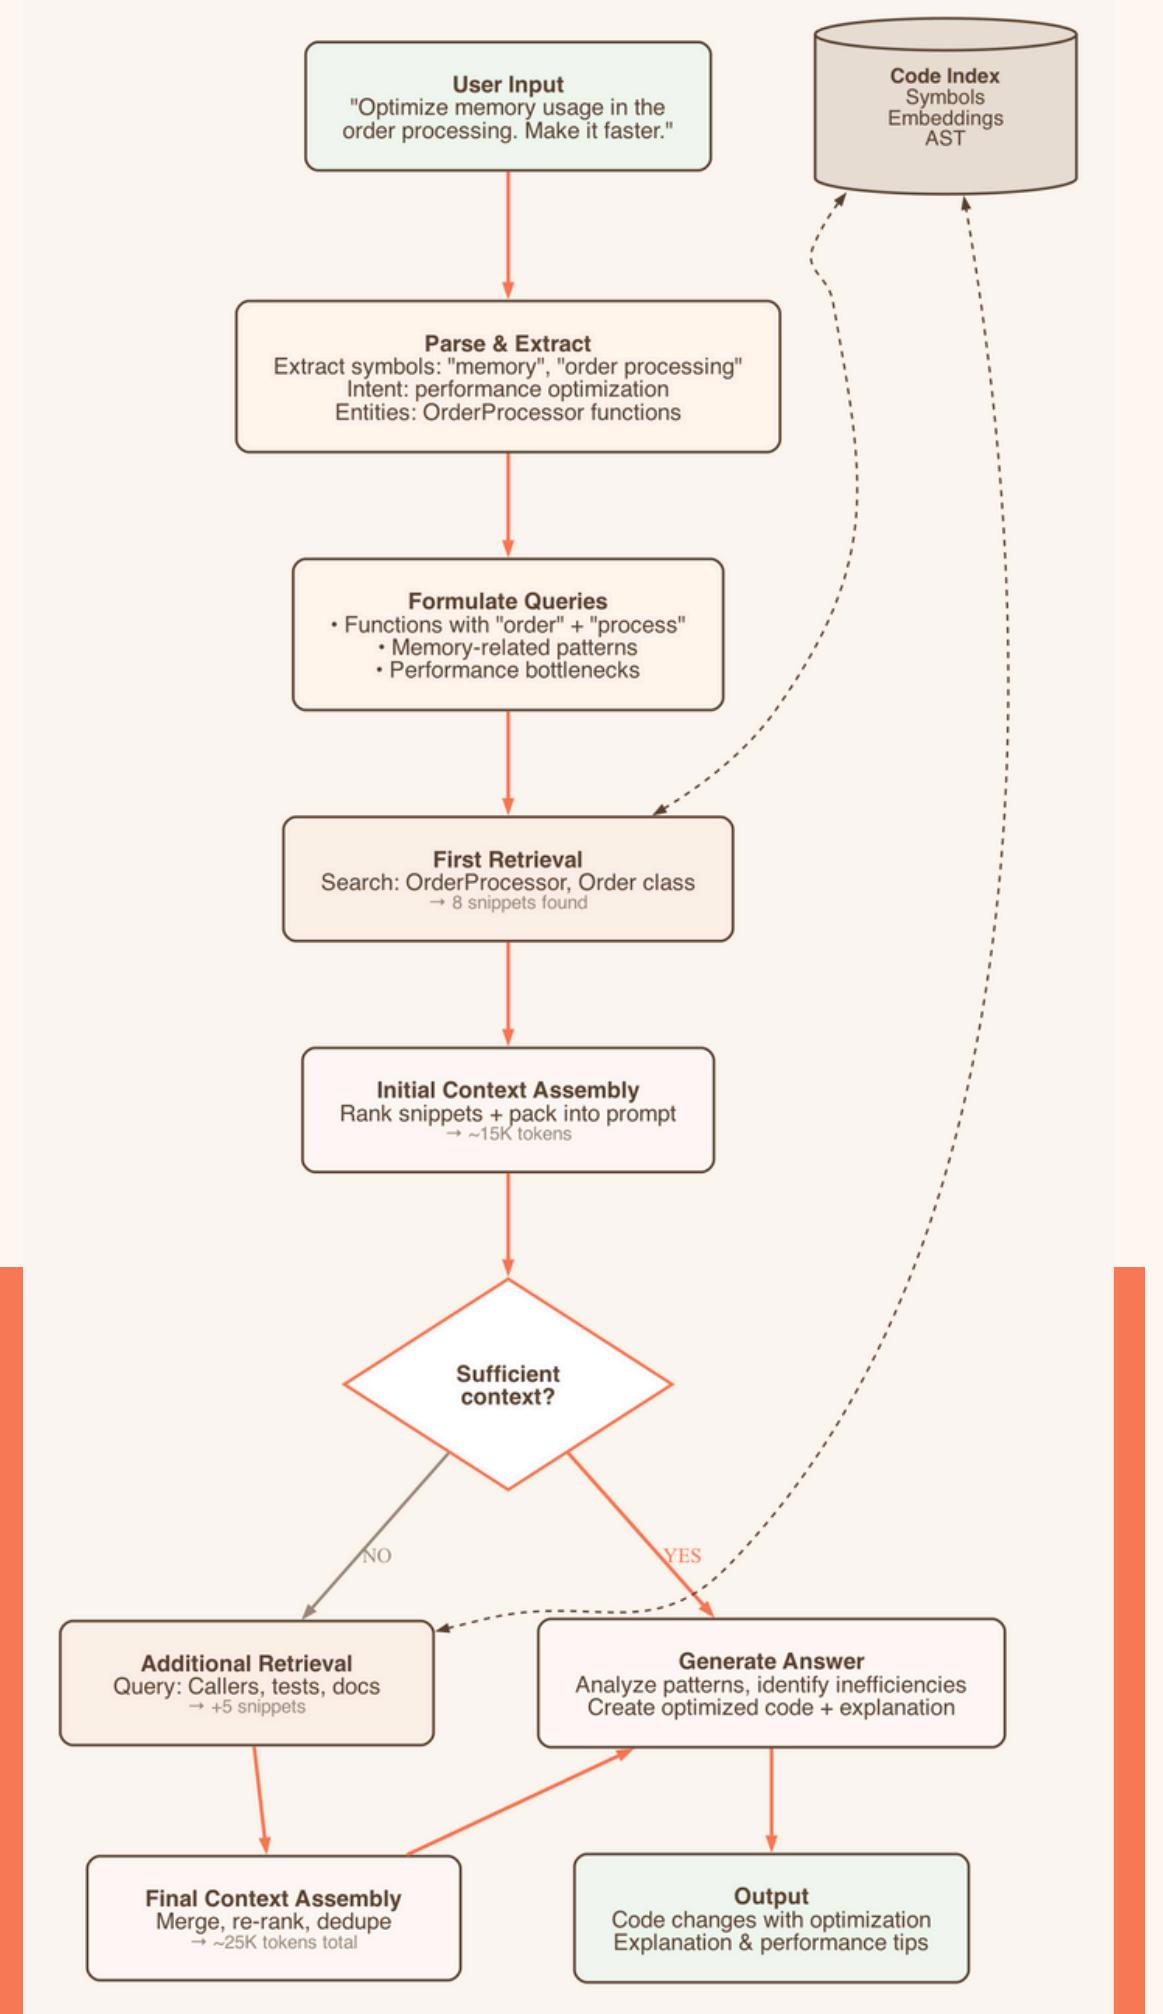
Determinism

sebagai salah satu referensi sebelum membuat strategi solusi

If search breaks or returns nothing,
the LLM can't "magic" the answer.



Example 1



Prompt:

"Optimize memory usage in the order processing. Make it faster."



Prompt parsing

Needs to guess intent, symbols, scope and acceptance criteria - pure vibe coding!



Context/search miss

Needs to find the right files and symbols in the workspace



Symbol resolution drift

Needs to figure out which processOrder to change, might be confused by dynamic inputs or name collisions



Tool execution failures (build/test/lint)

It needs to figure out which tool to execute and how to test that the task is implemented

Example 2: Prevent Vibe Coding!

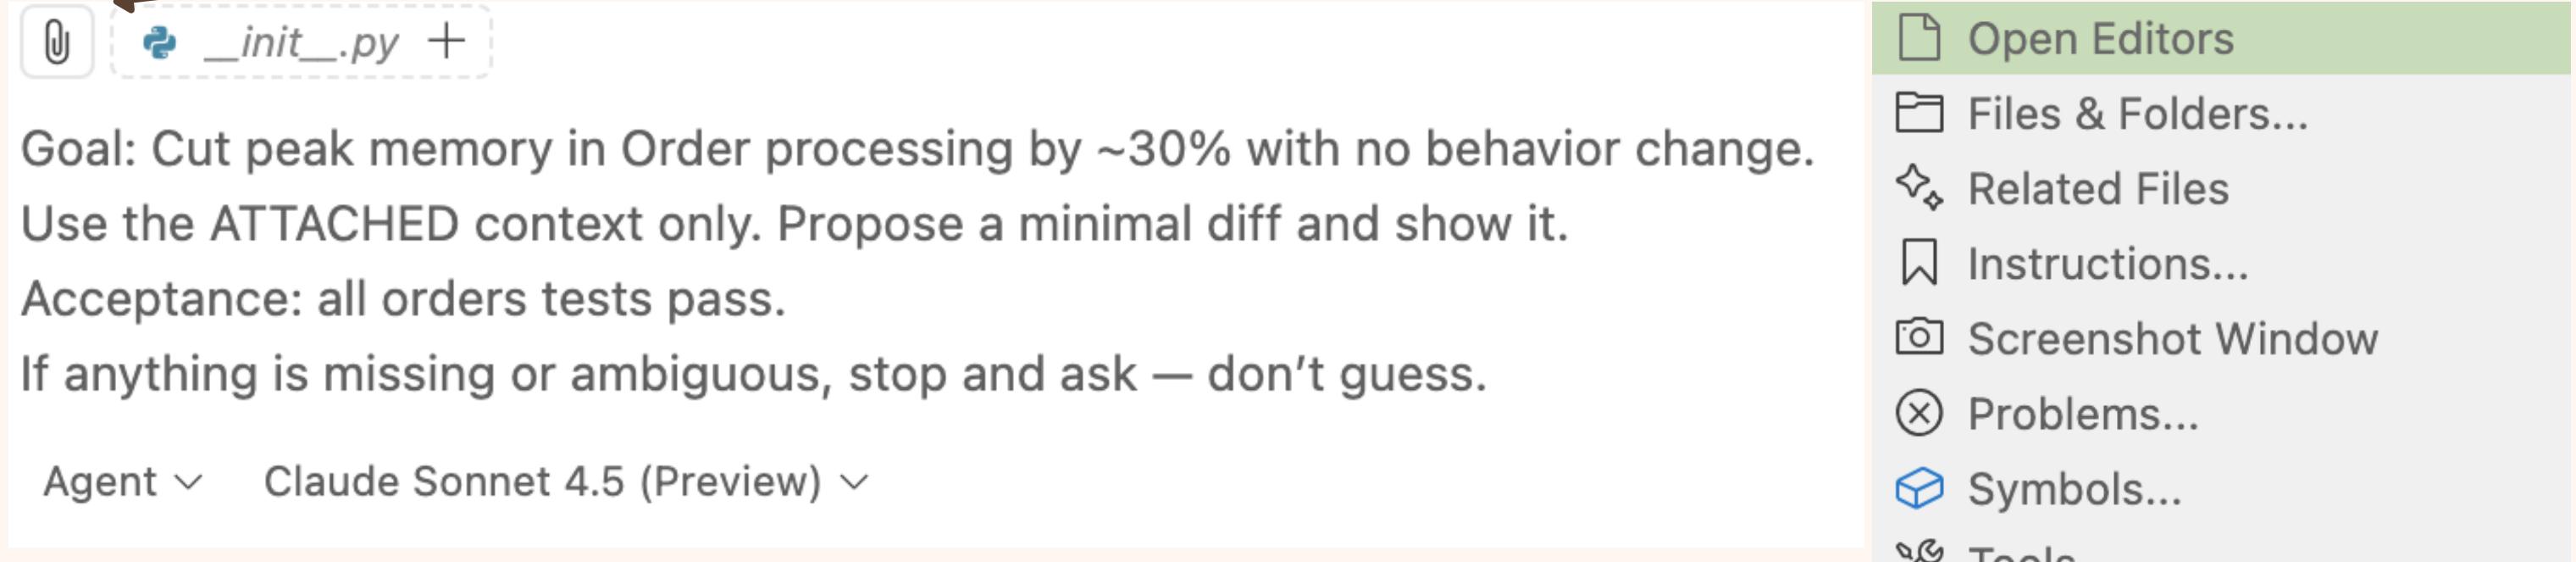
“Goal: Cut peak memory in Order processing by ~30% with no behavior change.
 Use the ATTACHED context only. Propose a minimal diff and show it.
 Acceptance: all orders tests pass.
 If anything is missing or ambiguous, stop and ask — don’t guess.”

Rule of Thumb:

- Avoid ambiguity
- State the goal
- Name acceptance
- Force model to ask and prevent guessing!
- Configure context such as tools, symbols, files etc



Add context



The screenshot shows a code editor interface with a Python file named `_init_.py`. Above the code area, there's a toolbar with icons for file operations. A context menu is open on the right side of the screen, listing several options: "Open Editors" (highlighted in green), "Files & Folders...", "Related Files", "Instructions...", "Screenshot Window", "Problems...", "Symbols...", and "Tools...". Two arrows point from the text above to the "Open Editors" item in the menu: one arrow points from the word "context" in the first sentence to the menu item, and another arrow points from the word "context" in the second sentence to the same menu item.

Goal: Cut peak memory in Order processing by ~30% with no behavior change.
 Use the ATTACHED context only. Propose a minimal diff and show it.
 Acceptance: all orders tests pass.
 If anything is missing or ambiguous, stop and ask — don’t guess.

Agent ▾ Claude Sonnet 4.5 (Preview) ▾

GUI Changes, Principals Stay The Same



Simplify search by adding context

Shorten the prompt; attach the facts (files, selections, symbols, docs).



Clean the context

Make sure the chat history is relevant, correct context is needed, cursor is at the right place



Prevent the LLM from guessing

Tell it what to do and when to ask.



Keep changes small; verify early

Prefer minimal diffs and short loops.



Don't overthink prompt engineering

Use a tiny template; let context carry the weight.

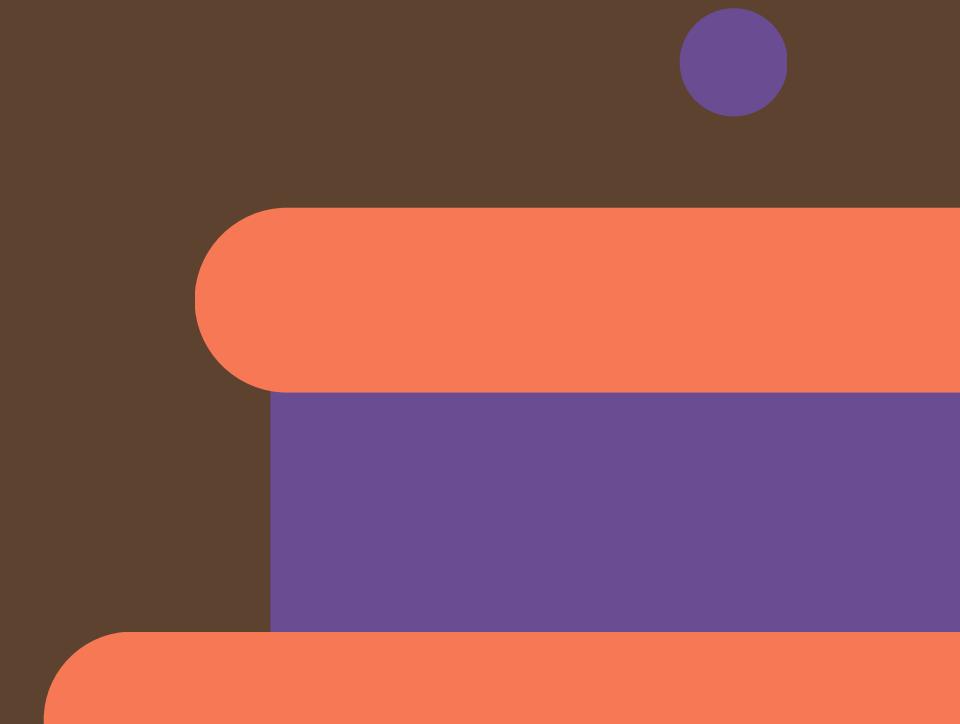


Acceptance gates

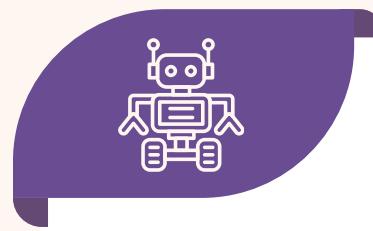
Encourage self-correction and self-reflection

Automation Paradigm & Tooling

Date: June 11, 2030

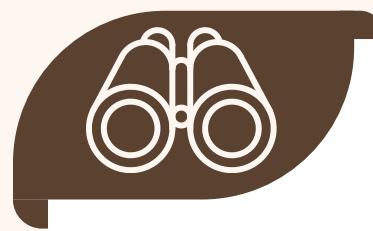


AI Changes The Paradigm: Machines Work, Humans Observer



Learn to delegate actions to AI

The more you successfully delegate the better



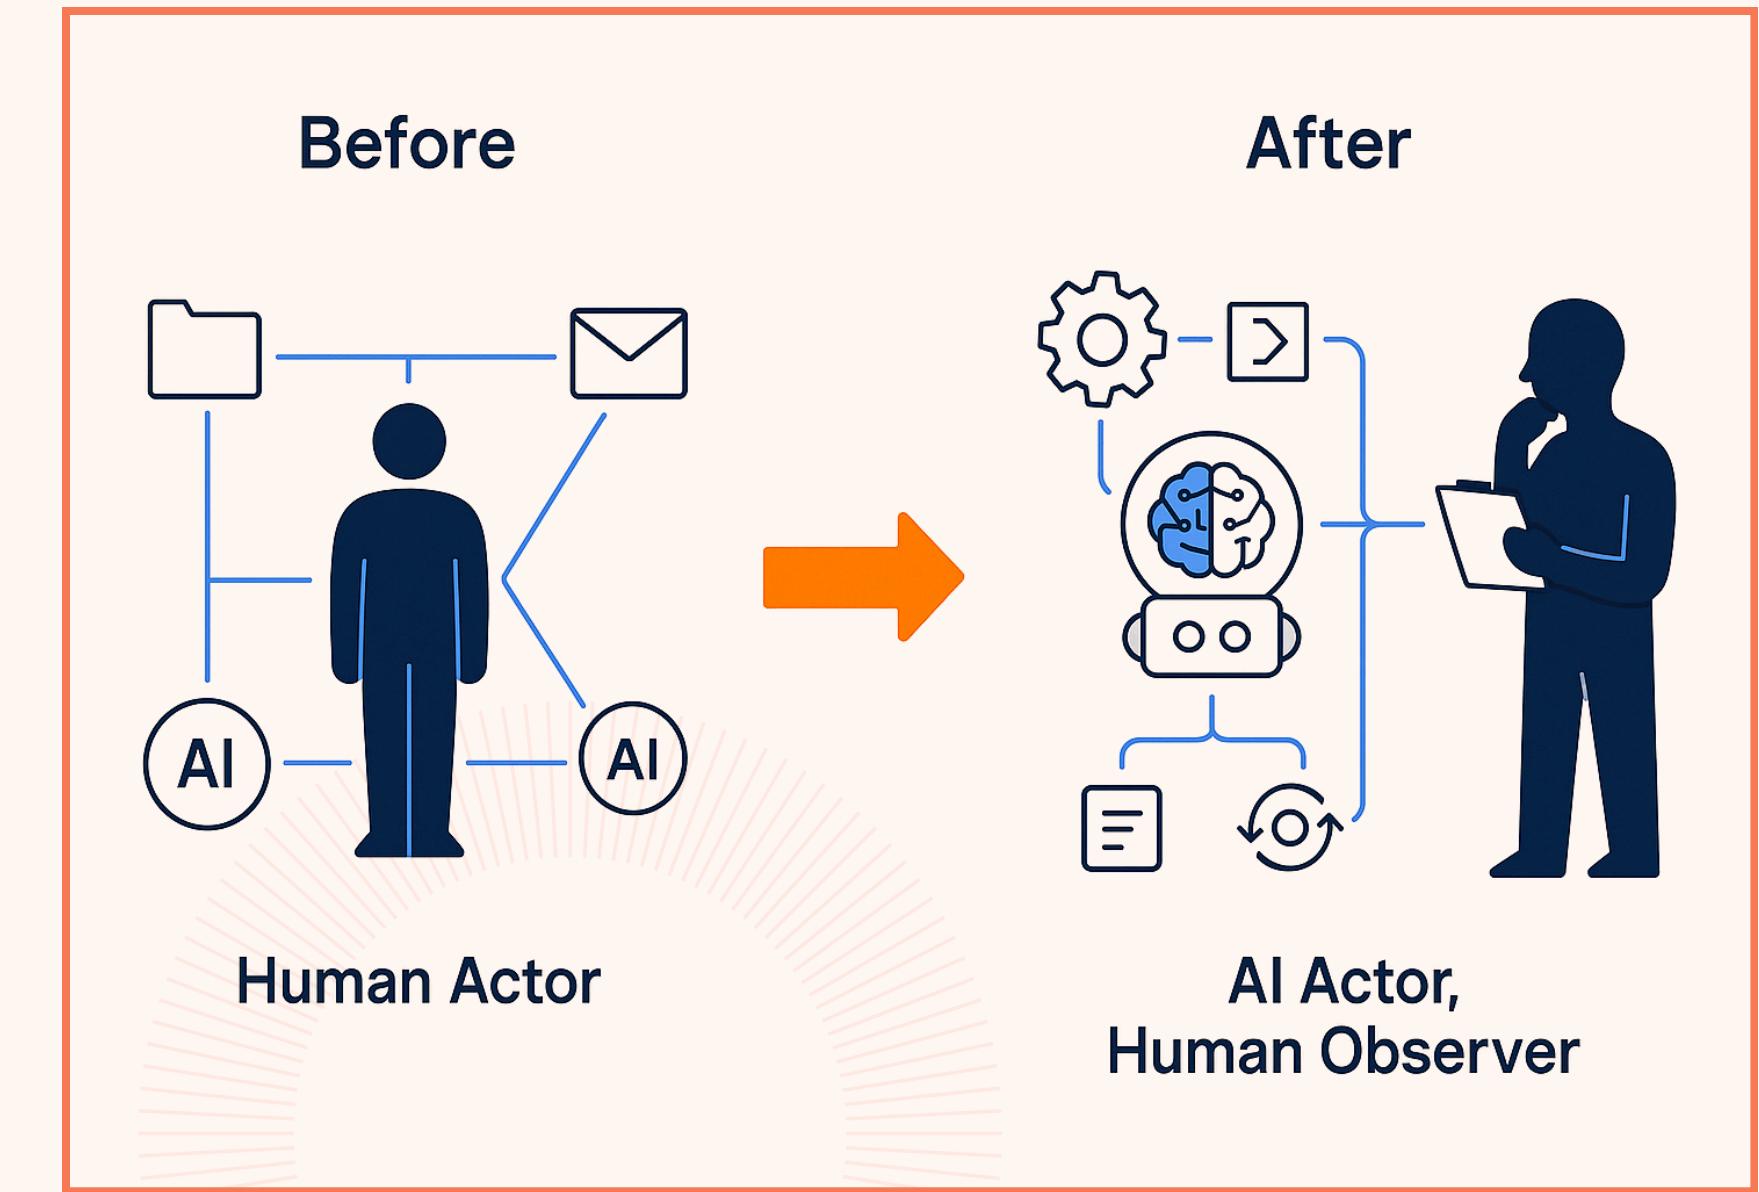
Learn to observe

Evaluate, consult, guide and ensure the quality



Learn to increase automation level

The goal is to move as far as possible in the degree of automation



L1



User as an
Operator

User directs and makes decisions, agent acts.

L2



User as a
Collaborator

User and agent collaboratively plan, delegate, and execute.

L3



User as a
Consultant

Agent takes lead but consults user for expertise/preferences.

L4



User as an
Approver

Agent engages user only in risky or pre-specified scenarios.

L5



User as an
Observer

Agent operates with full autonomy under user monitoring.

User Involvement

Agent Autonomy

Will We All Get Fired?

**GitLab**

Bill Staples
CEO of GitLab

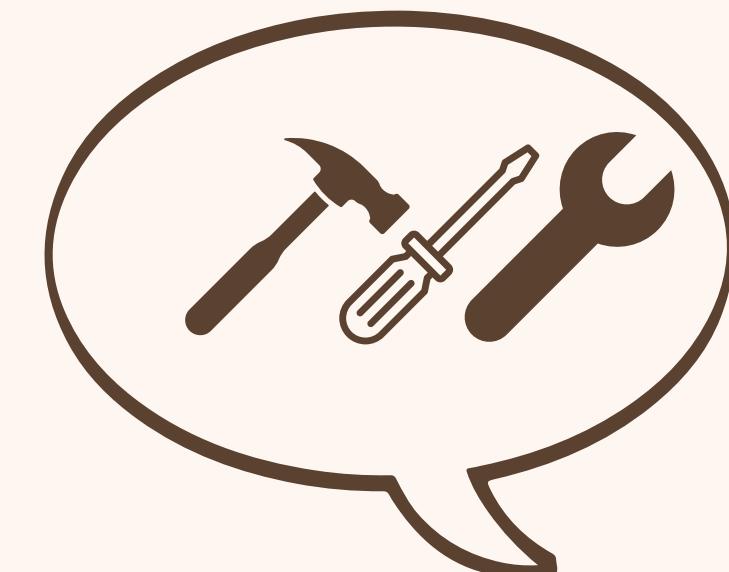
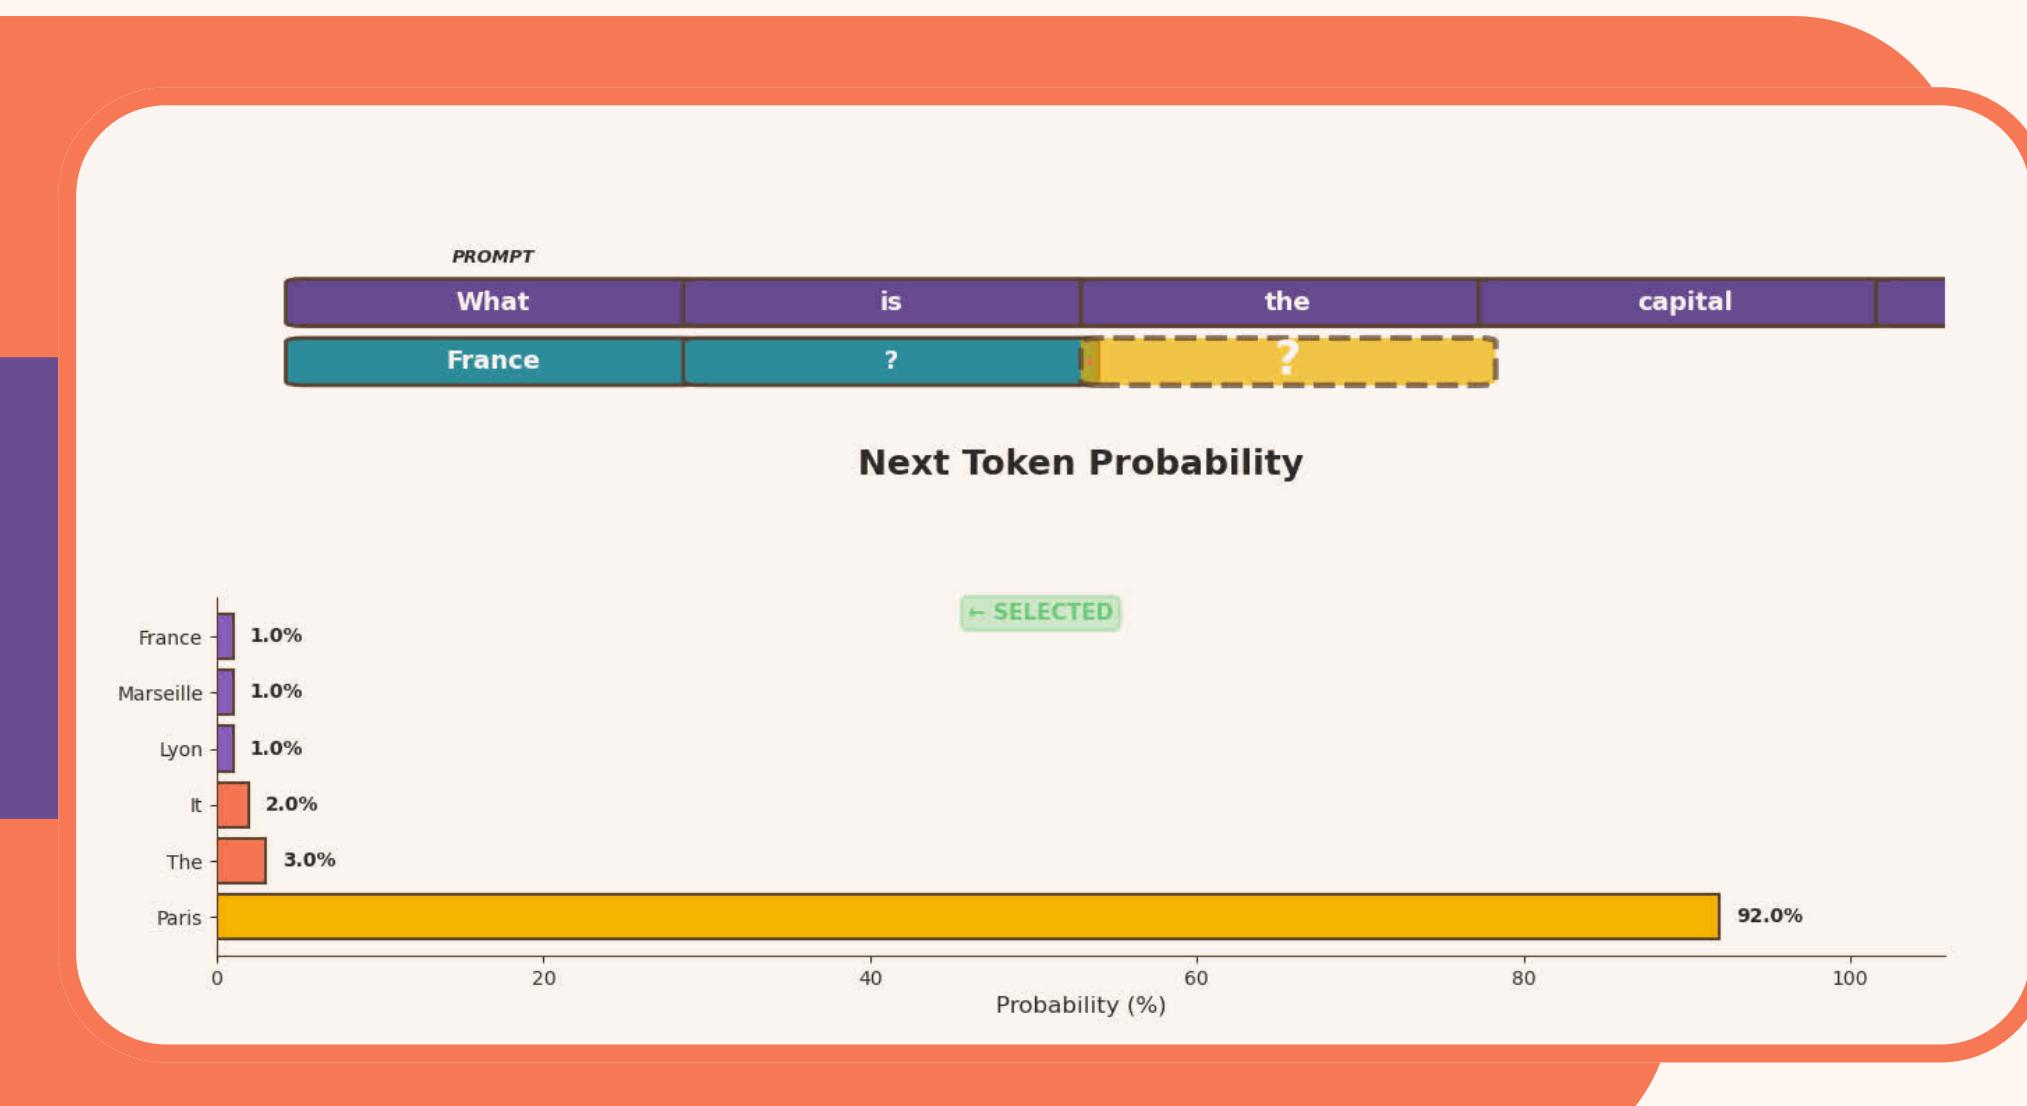
“

“Imagine what can happen when we double, triple, quadruple, 10x the number of people who can now create software through natural language.”

”

Coding LLMs: Tool Usage Prediction

The latest models starting from o1 series, Claude 3.5, Gemini 2.0
are explicitly trained for tool selection

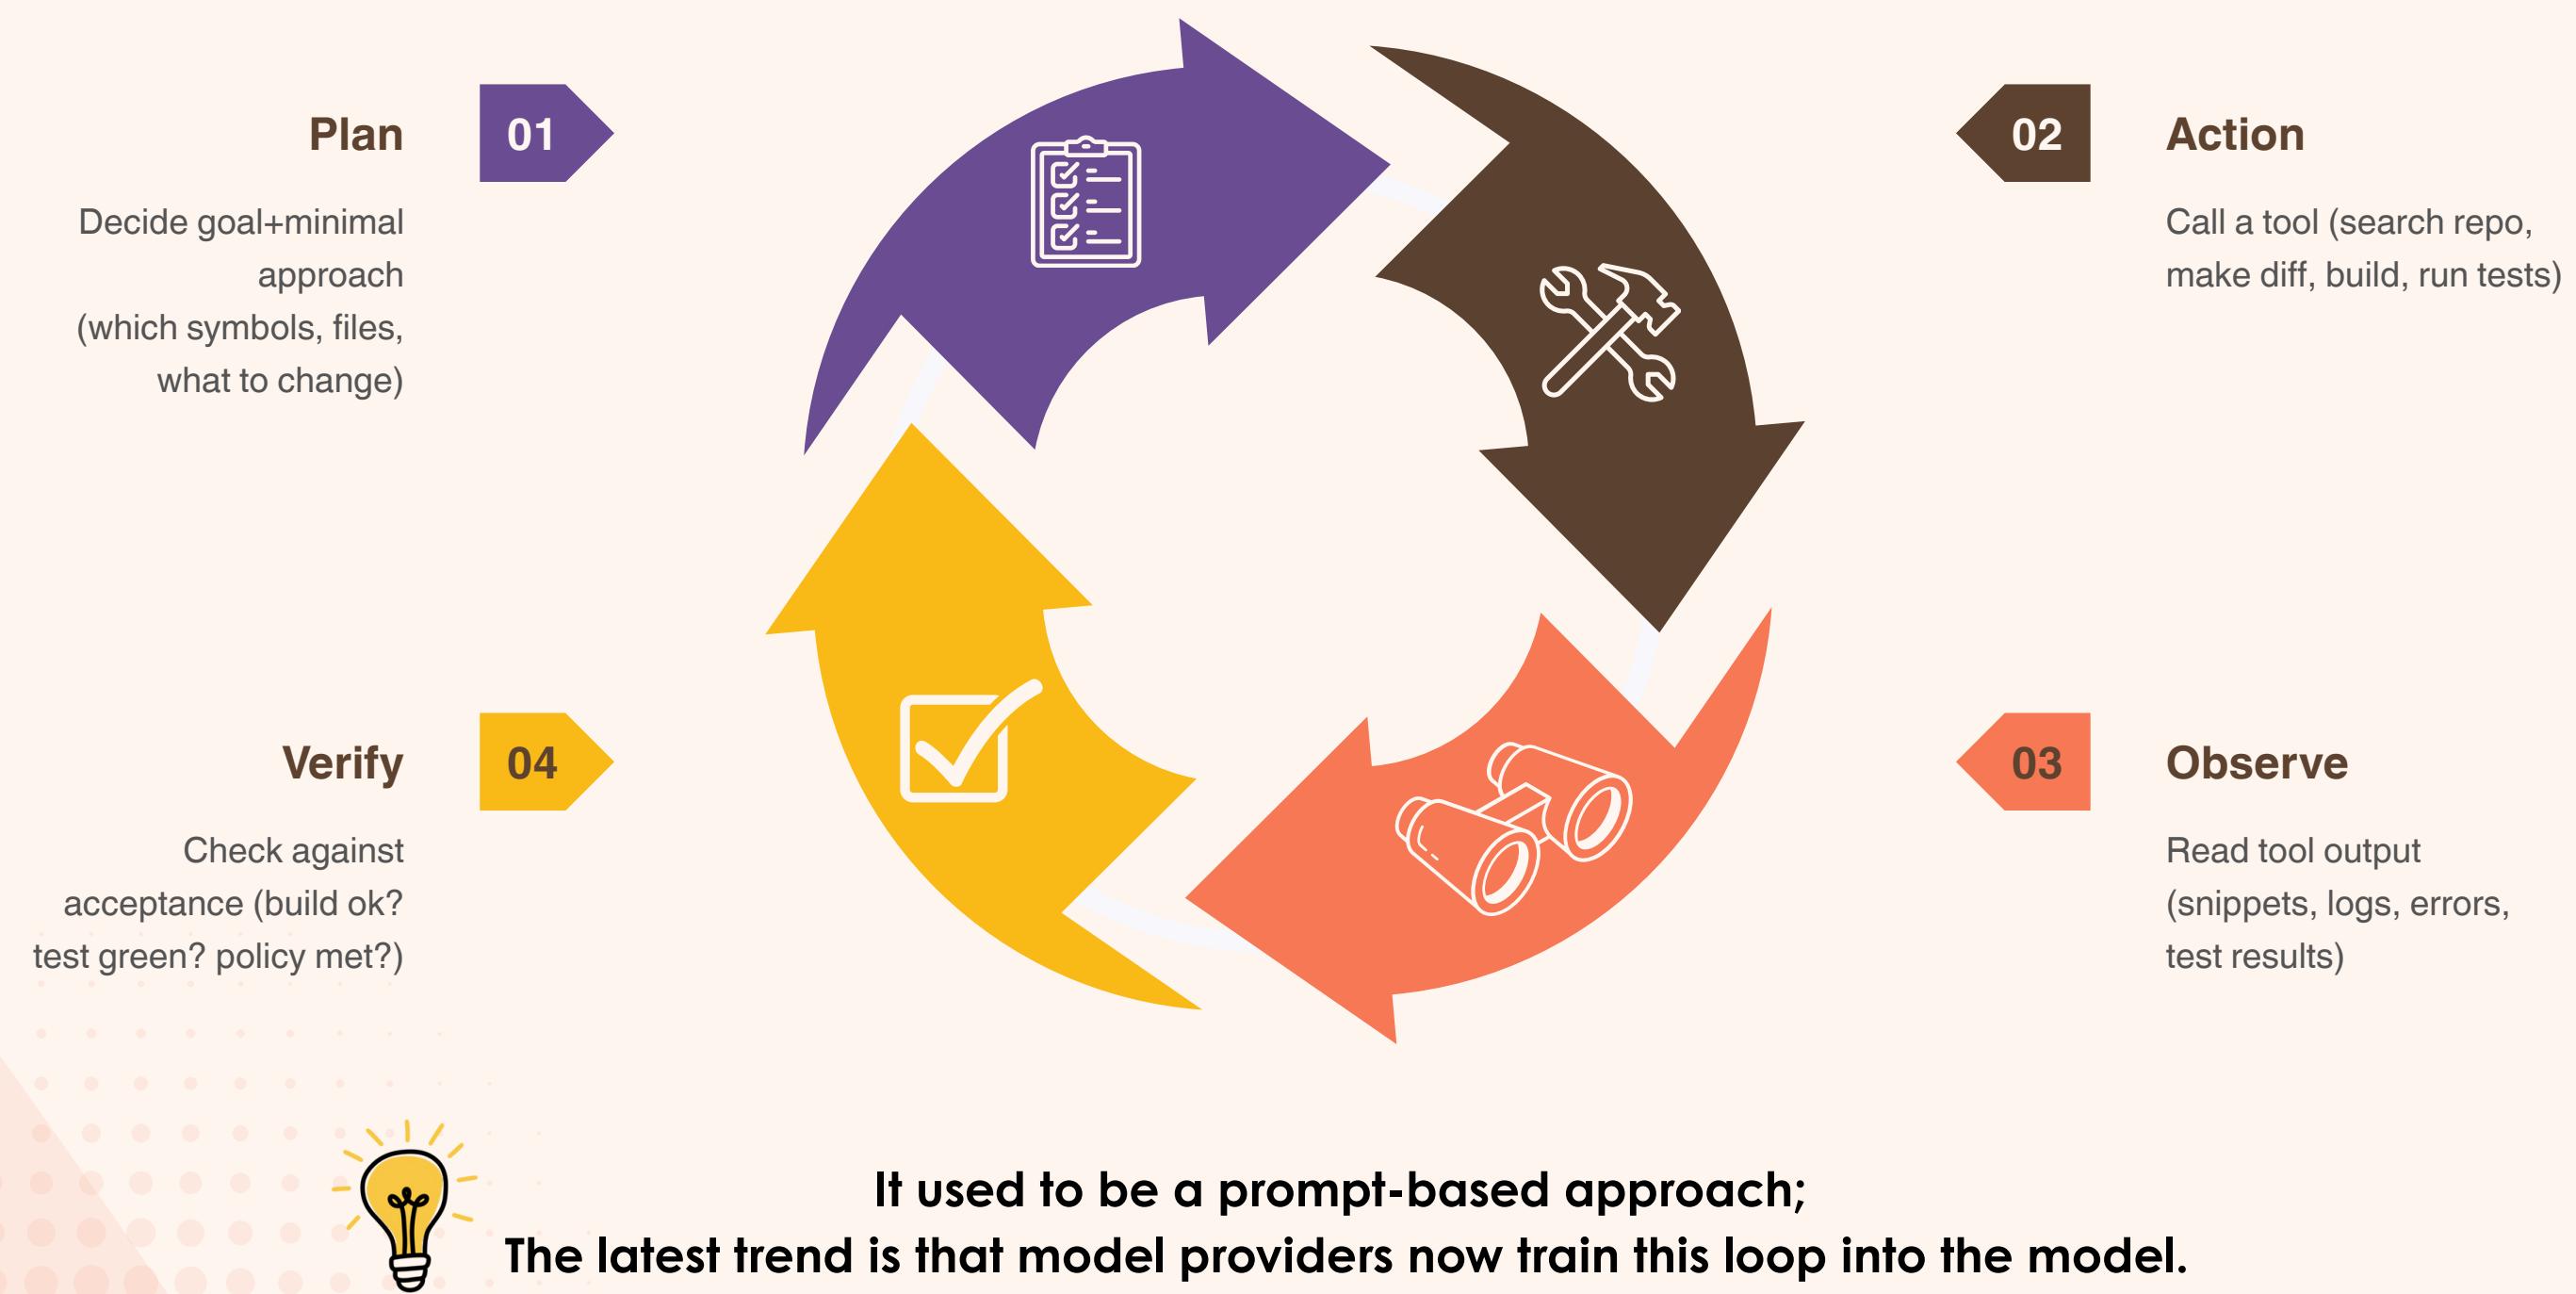


Modern LLMs have two jobs:
predict next word and tools



Thought-Action-Observation Cycle

How AI Agents are Defined and Trained



What Training Data Look Like

```

"issue_title": "Rounding error: 19.995 becomes 19.00 instead of 20.00",
"initial_tests": "failing",
"trajectory": [
{
  "step": 1,
  "think": "Reproduce the failure to see the exact assertion and stack trace.",
  "act": { "tool": "pytest", "args": ["-q"] },
  "observe": {
    "exit_code": 1,
    "stdout": "F..F...\nE AssertionError: expected 20.00, got 19.00\nFAILED tests/test_discounts.py::TestDiscounts::test_rounding[19.995]\n",
    "stderr": ""
  }
},
{
  "step": 2,
  "think": "The failure is in pricing. Open the suspected module to inspect rounding."
  "act": { "tool": "open_file", "args": ["discounts/pricing.py"] },
  "observe": {
    "content_head": "def round_price(x: float) -> int:\n      # TODO: simplify rounding"
    "lines": 120
  }
},
{
  "step": 3,
  "think": "int(x) truncates instead of rounding. Replace with round-half-up behavior"
}
]

```

The agent is trained to generate trajectories:

- think
- act
- observe

The more context it gets that fit this pattern e.g. it is told how to act, which context to look at, what to observe, the less iterations it does

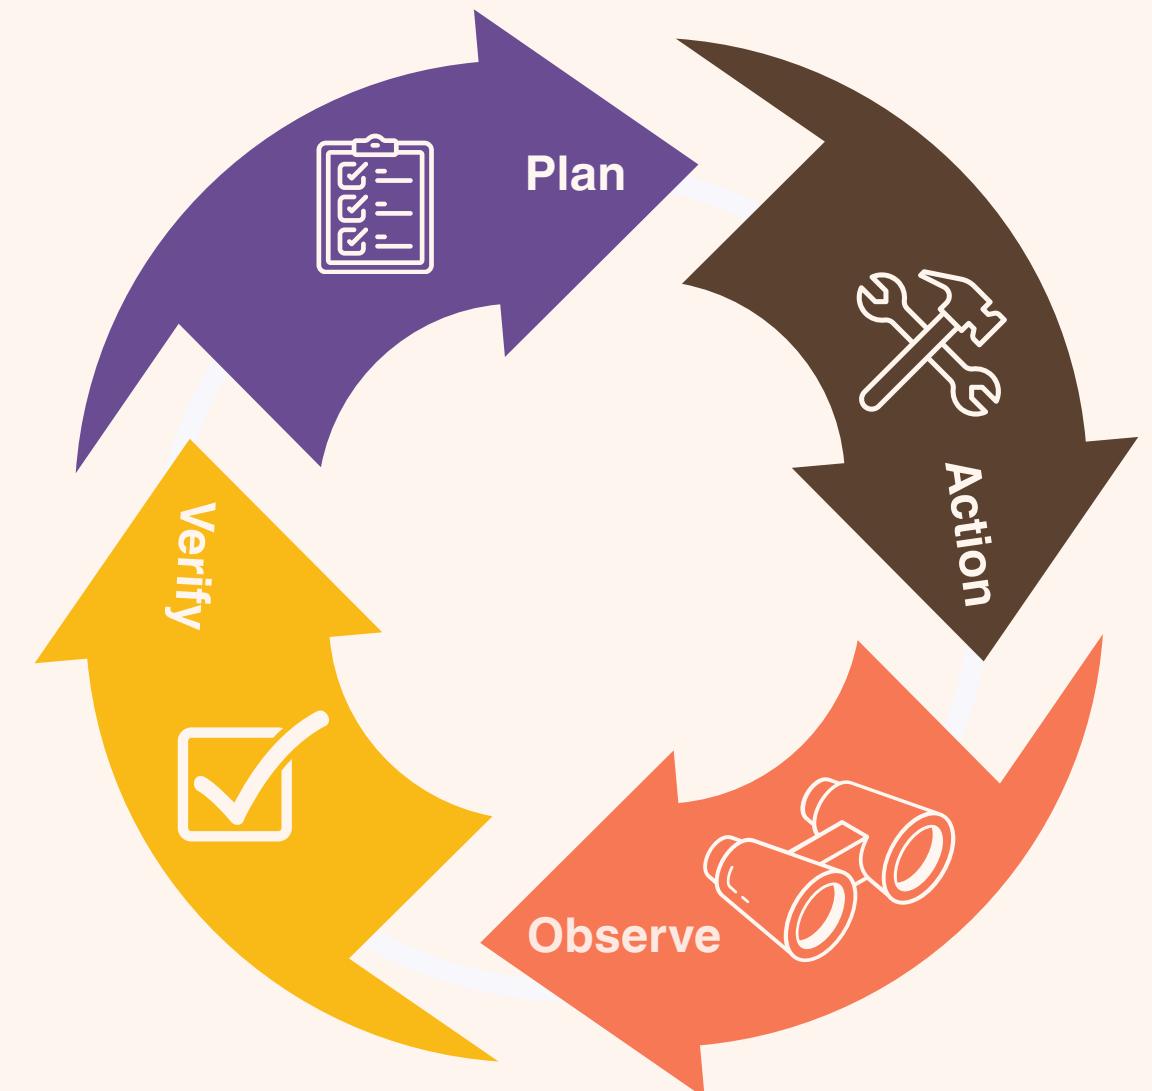
Coding Agents

Not just a theory — this is how modern coding models are trained/aligned

AI coding agent

A program that uses a large language model (LLM) to plan, act, observe, and verify when making changes to software:

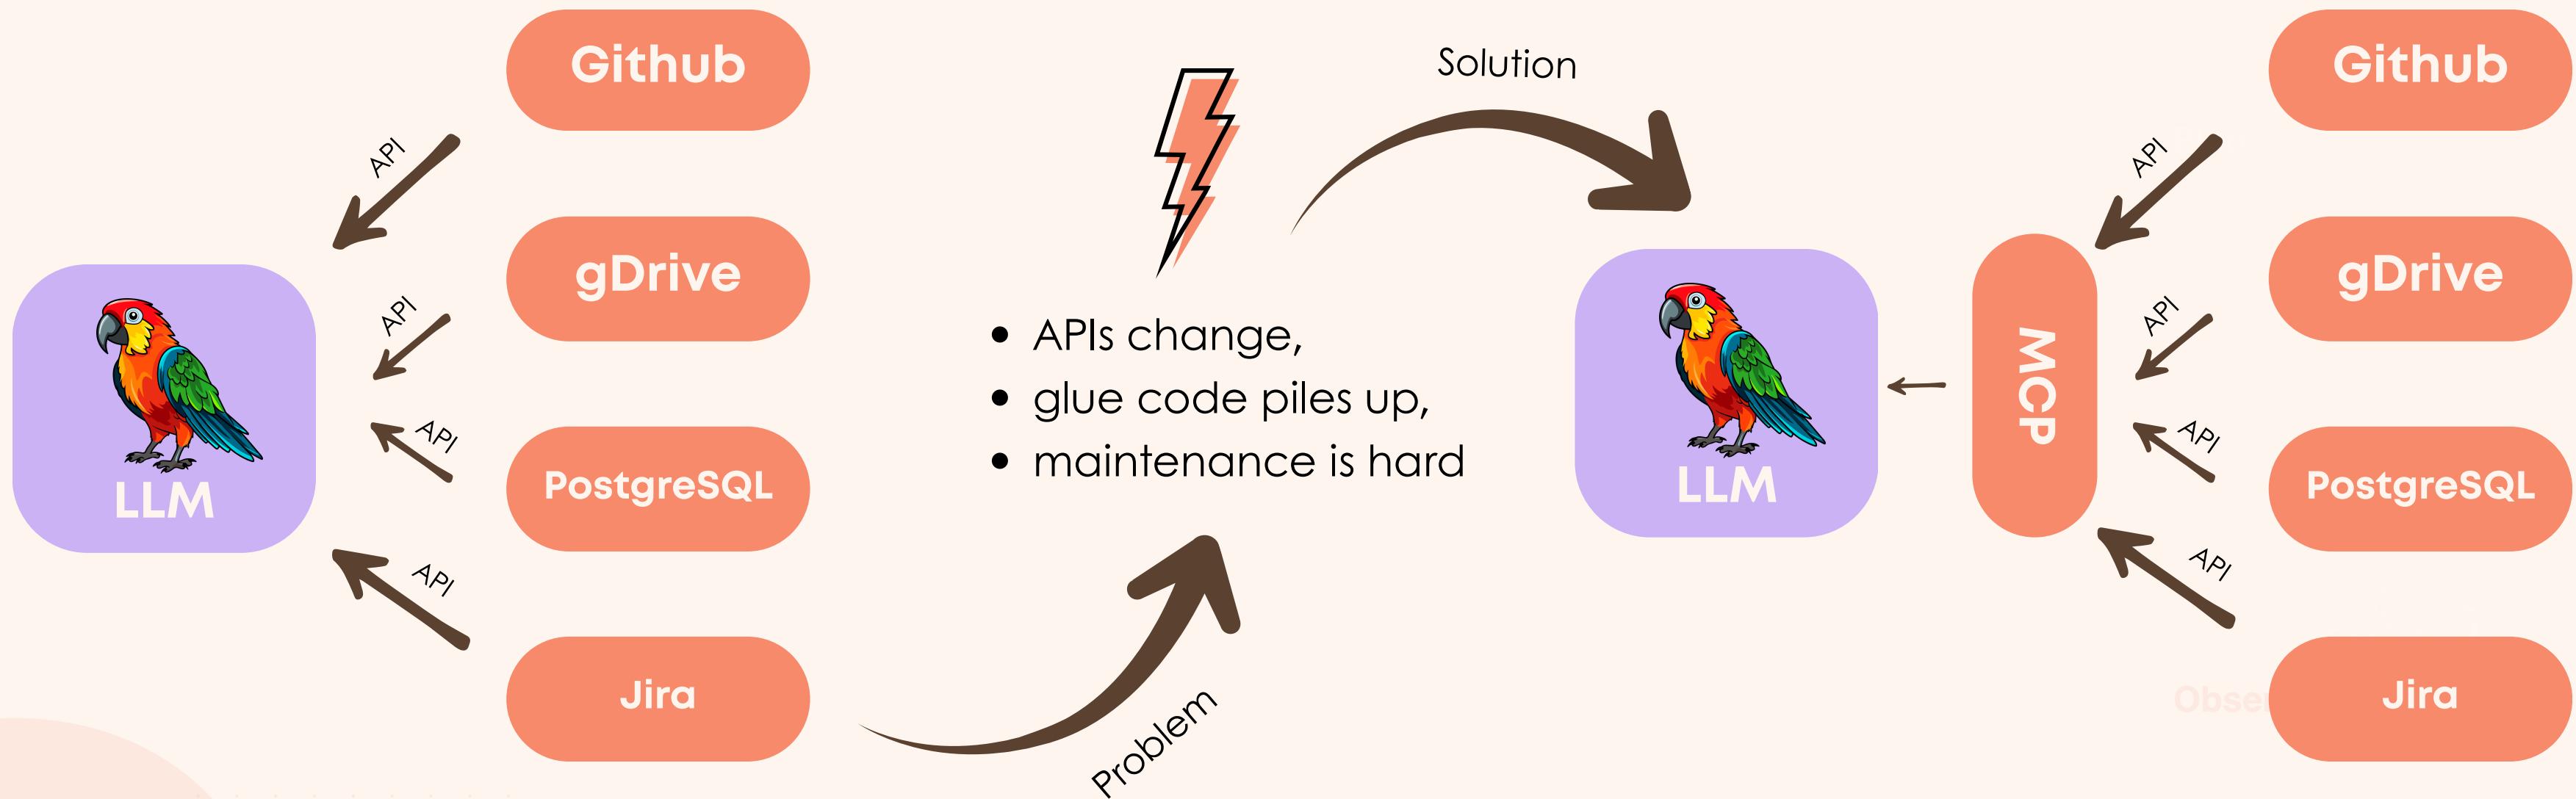
- **Plan**: derives the goal, outlines steps, and sets acceptance criteria
(e.g., tests must pass, performance target).
- **Act**: chooses and invokes concrete tools
(e.g. repo search, edit/patch, compiler, test runner).
- **Observe**: reads and interprets the environment's feedback
(e.g. tool outputs, logs, errors, code diffs, runtime results).
- **Verify**: checks results against the acceptance criteria
(e.g build/test pass, style/policy met, perf budgets)



Behind the scenes, **every modern model runs a plan → act → observe → verify cycle.**
Reduce guessing: attach precise context and state clear acceptance tests.

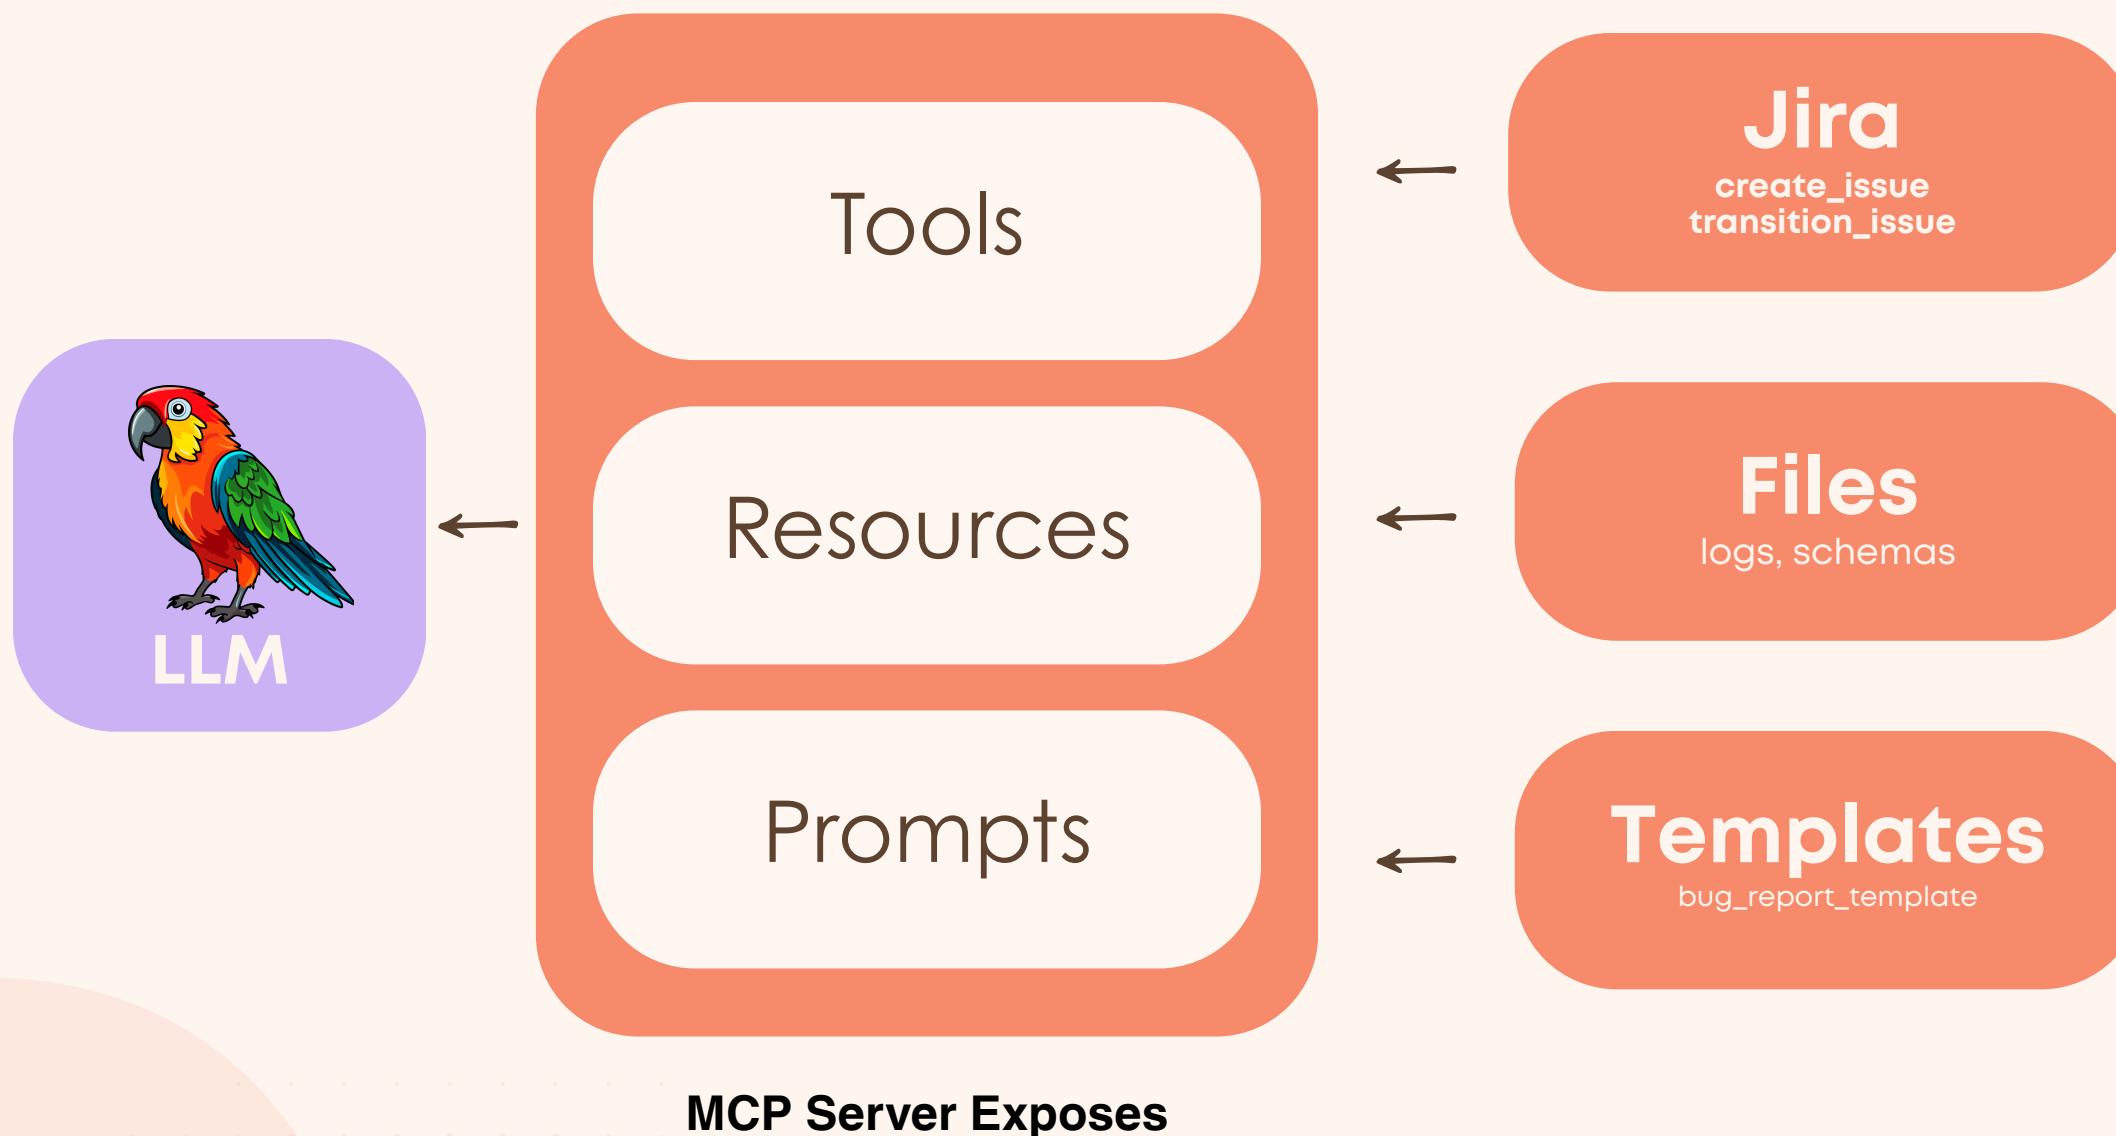
MCP Servers

Unified way to bring context and tools to LLMs



MCP Servers

Unified way to bring context and tools to LLMs



GitHub Copilot only supports tools

Workaround:

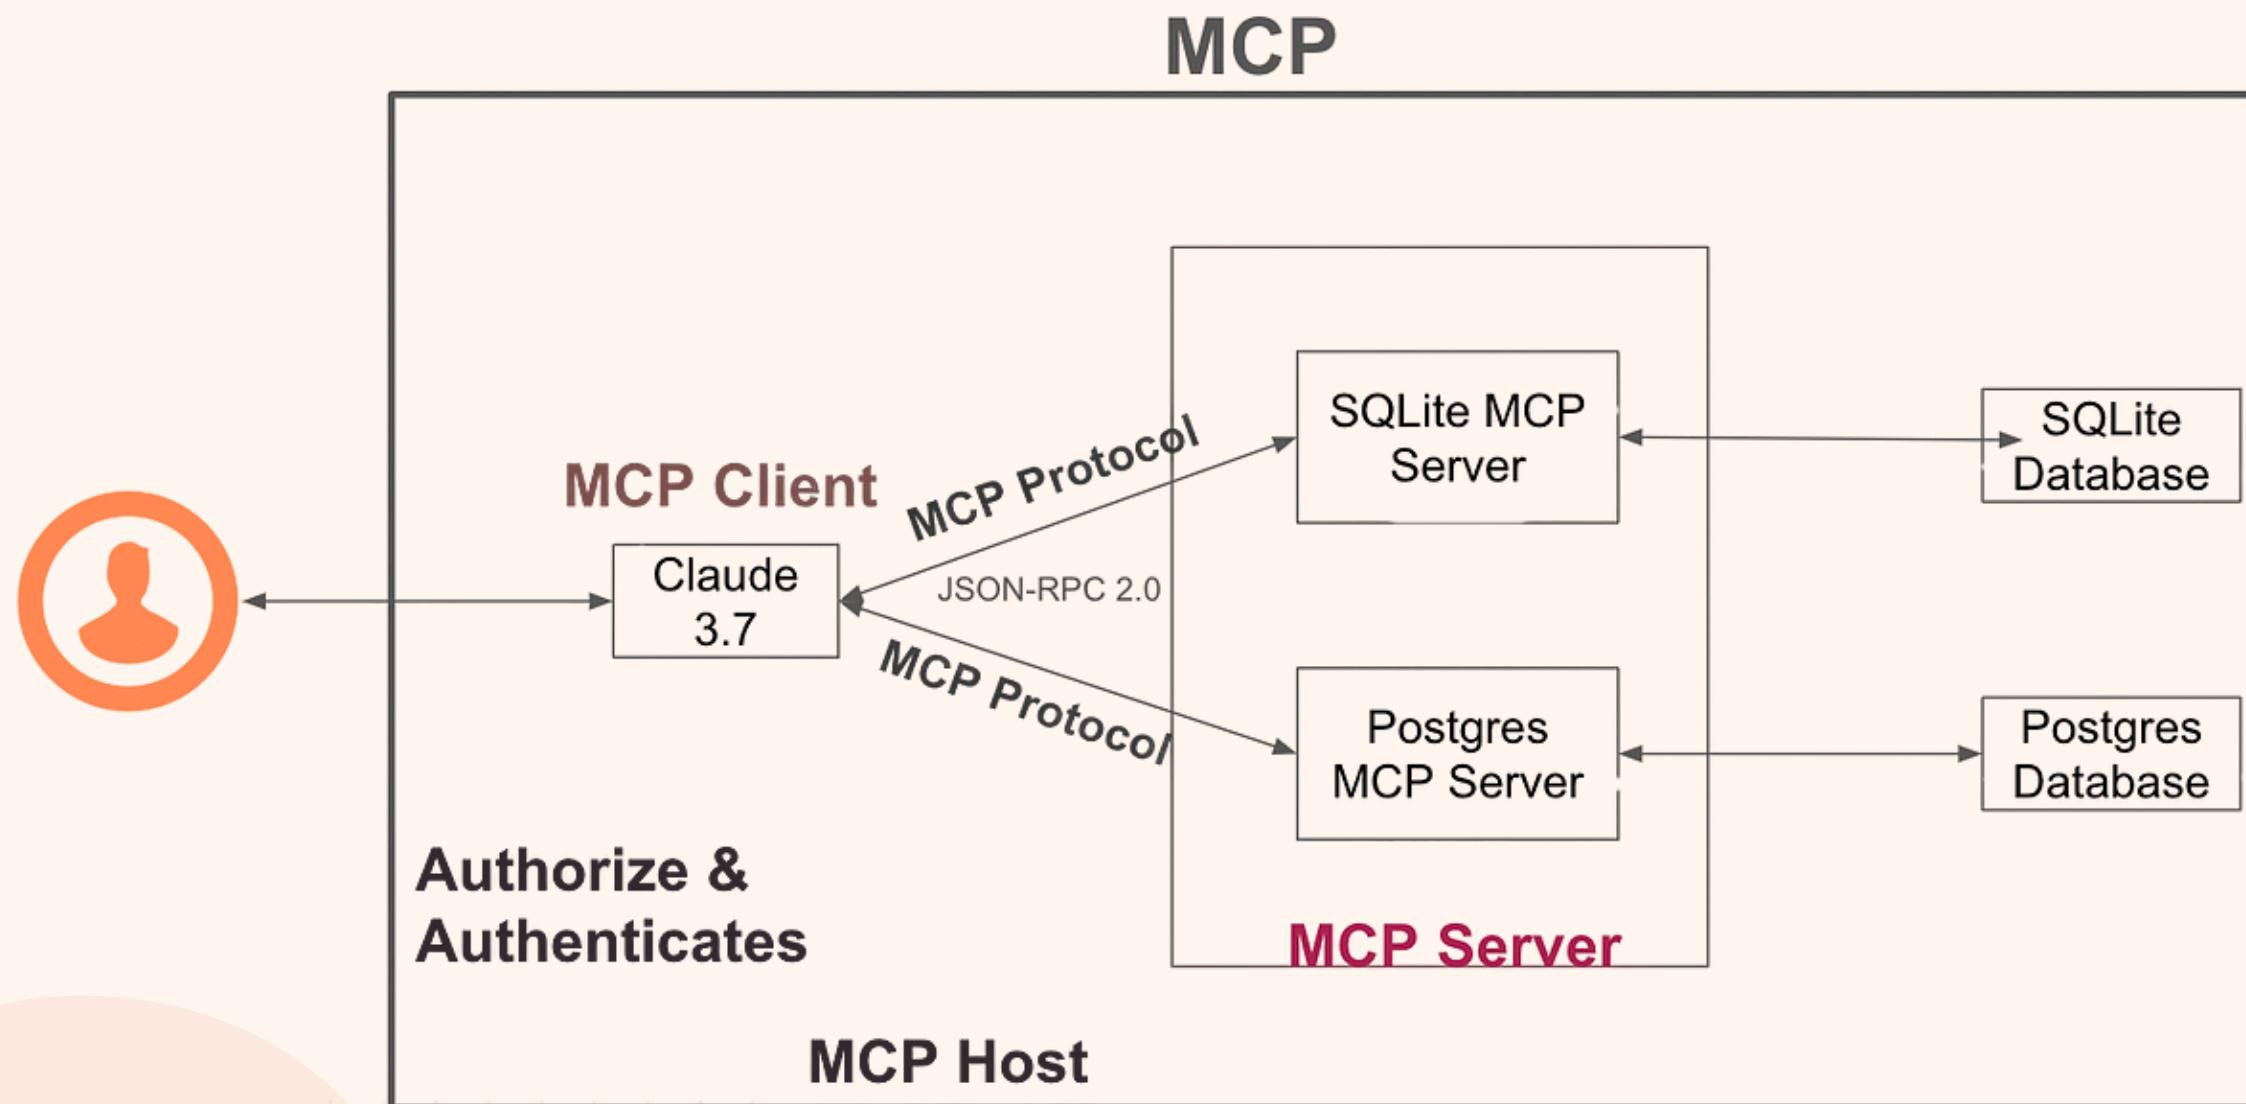
- wrap read-only context as tools (e.g., read_logs, fetch_schema)
- wrap reusable prompt templates as a tool (e.g., get_prompt_template({name, vars})).

Best Practices:

- **In-repo:** Add context (files, selections, symbols) with Copilot's native controls.
- **Off-repo:** Use MCP tools to fetch external schemas, logs, APIs, etc.

MCP Servers

Unified way to bring context and tools to LLMs



JSON-RPC 2.0

Remote Procedure Call protocol using JSON

- Every custom MCP server **MUST** implement a JSON-RPC endpoint that Copilot can call.

JSON-RPC Methods

initialize
handshake
with client

tools/list
list available
tools

tools/call
Execute a
tool

shutdown
Clean
shutdown

Must

Should

MCP Servers

Unified way to bring context and tools to LLMs

```

@app.post("/mcp")
async def mcp(request: Request):
    payload = await request.json()
    method = payload.get("method")
    rpc_id = payload.get("id")

    if method == "initialize":
        return {"jsonrpc": "2.0", "id": rpc_id, "result": {
            "protocolVersion": "2024-11-05",
            "capabilities": {"tools": {}},
            "serverInfo": {"name": "my-server", "version": "1.0"}
        }}

    if method == "tools/list":
        return {"jsonrpc": "2.0", "id": rpc_id, "result": {
            "tools": [
                {
                    "name": "my_custom_tool",
                    "description": "Does something amazing",
                    "input_schema": {"type": "object", "properties": {}}
                }
            ]
        }}

    if method == "tools/call":
        tool_name = payload["params"]["name"]
        args = payload["params"]["arguments"]

        if tool_name == "my_custom_tool":
            # Your business logic here
            result = do_something_amazing(args)
            return {"jsonrpc": "2.0", "id": rpc_id, "result": {
                "content": [{"type": "text", "text": result}]
            }}

    return {"jsonrpc": "2.0", "id": rpc_id, "error": {
        "code": -32601, "message": "Method not found"
   }}

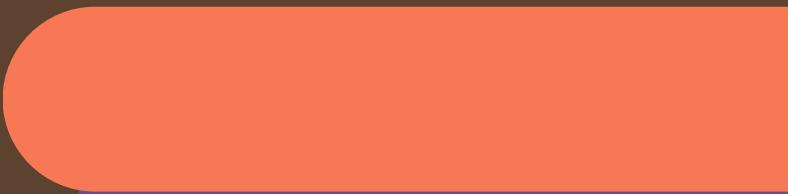
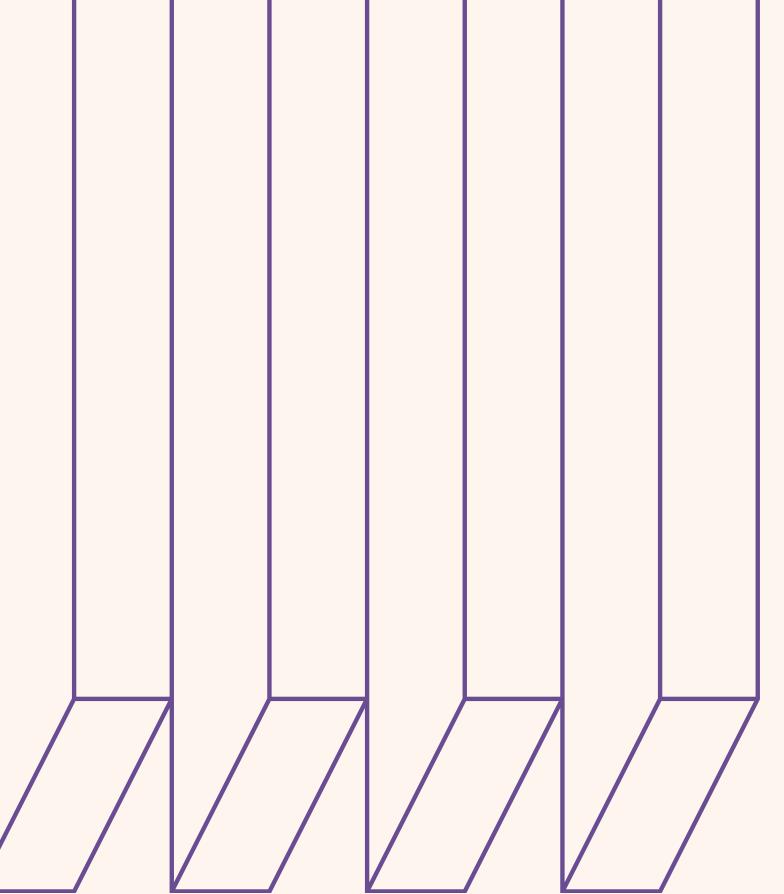
```

Implementation Checklist

- JSON-RPC endpoint (POST /mcp)
- Handle **initialise** method
- Handle **tools/list** method
- Handle **tools/call** method
- Define your custom tools
- Implement tool business logic (what it actually does)
- Return proper JSON-RPC responses
- Configure VS Code MCP settings

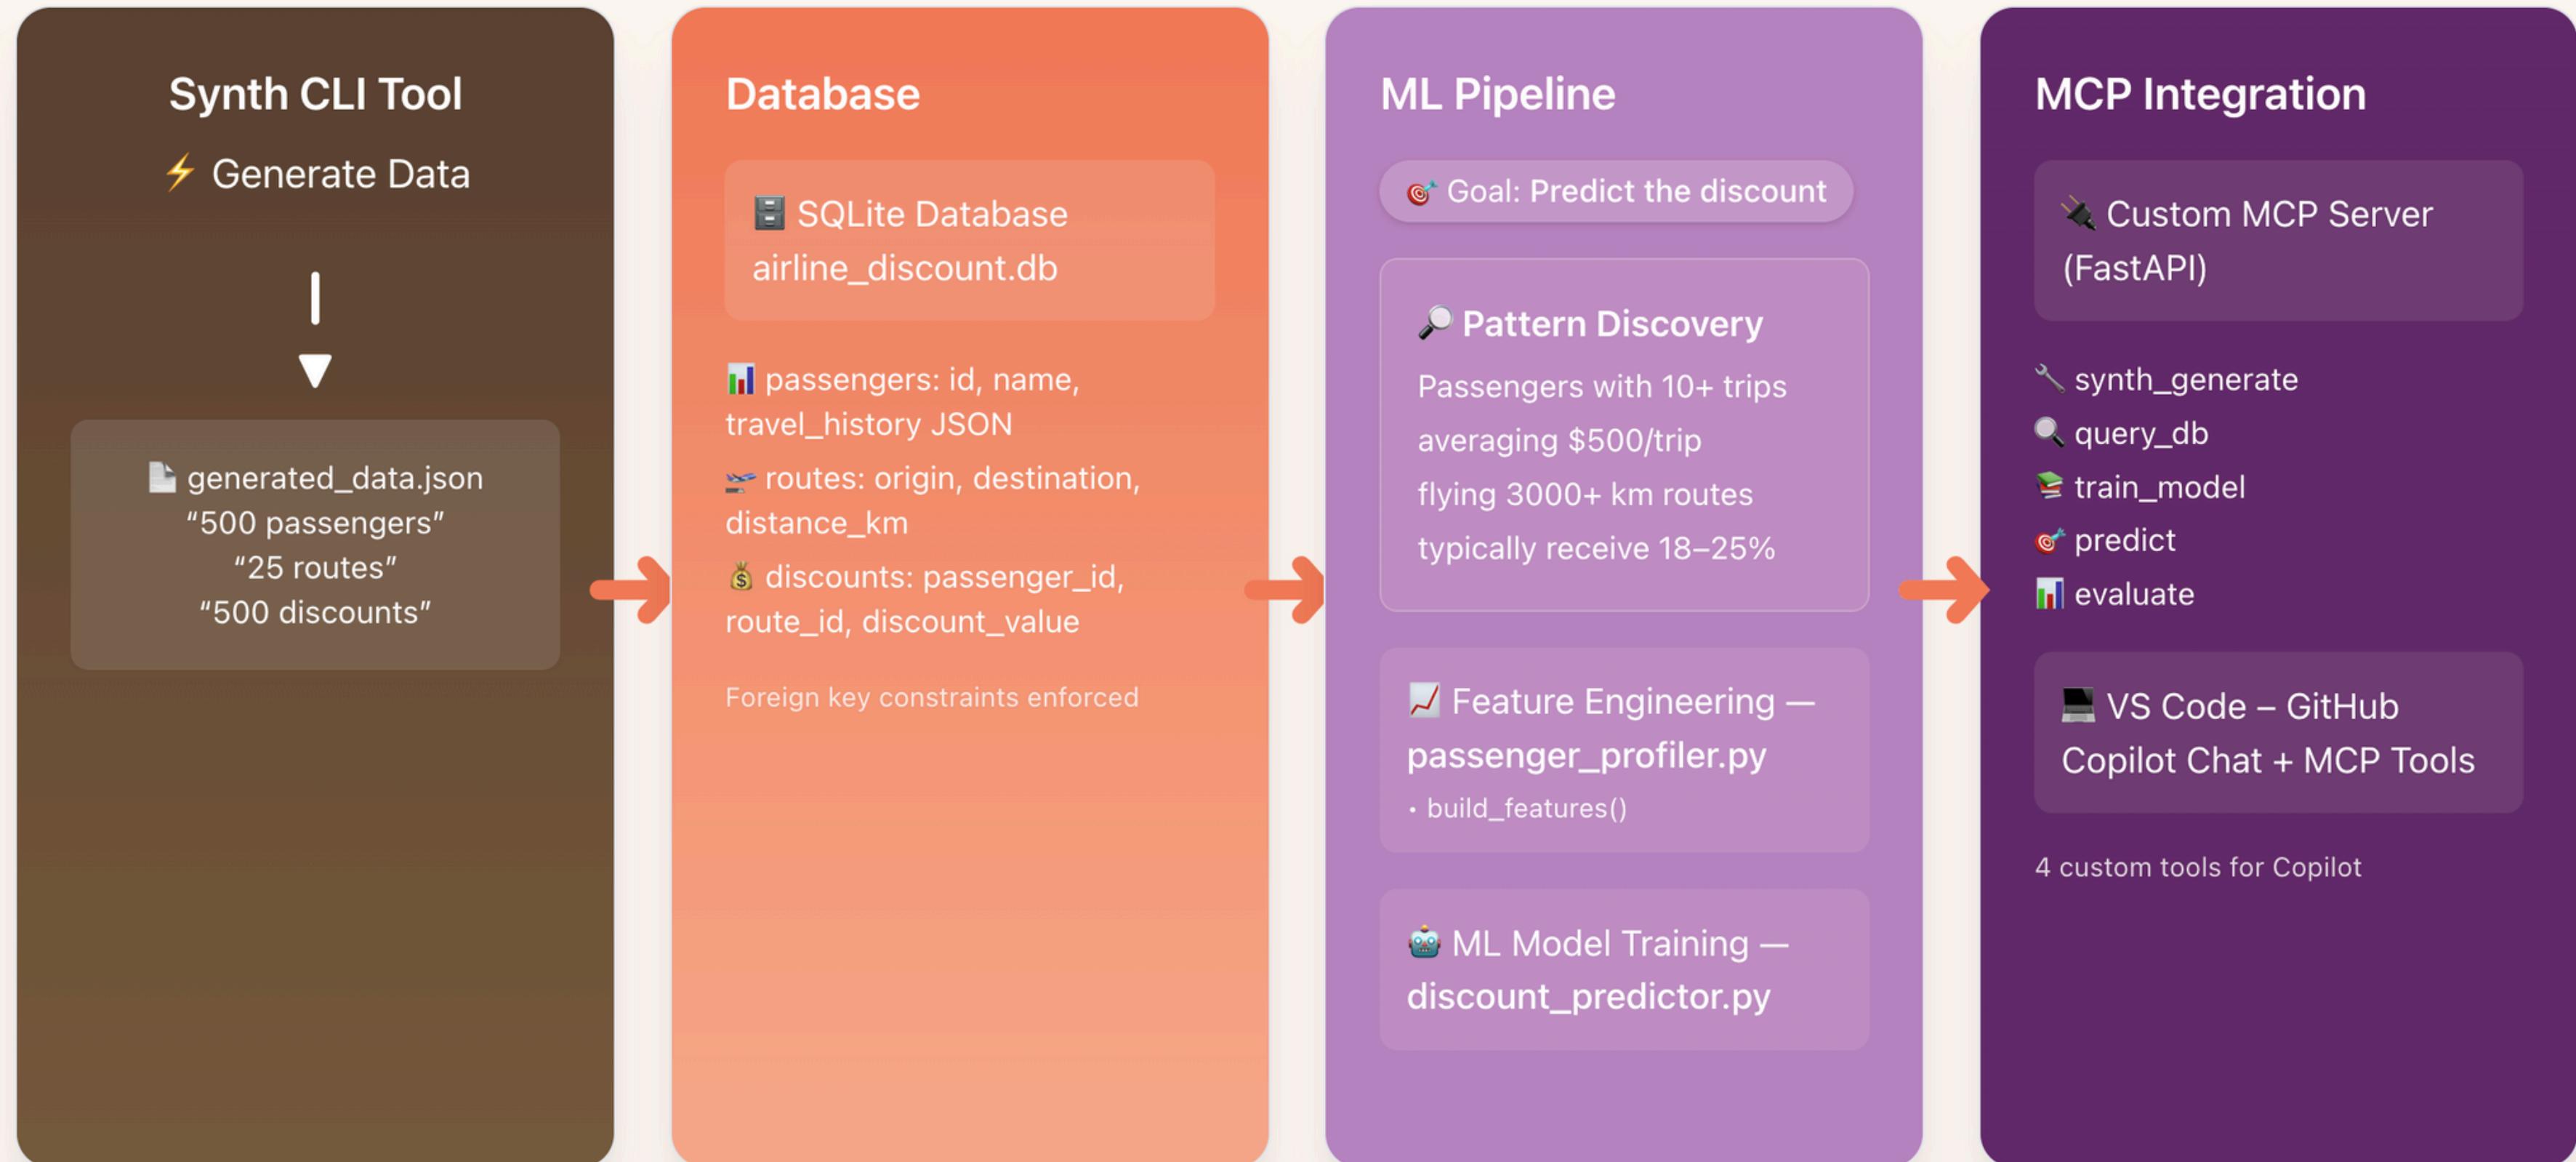
Practical Lab

Date: June 11, 2030



AIRLINE DISCOUNT ML PROJECT

Check VENV FIRST!!!!!!!!!



Airline Discount ML Project

Predicts airline discount percentages for passengers based on travel history and route data

Uses synthetic data only (SQLite database) – no real customer information

Implements a complete ML pipeline: data → features → model → predictions

1

2

3

EXERCISE ONE

SETUP

Set up GitHub Copilot Instructions

Install repo-wide and path-scoped

Enforce "ask, don't guess" behavior

Configure Copilot to request clarification

Use path-scoped instructions with applyTo globs

Apply different coding rules to specific parts of your codebase

Verify instruction loading in VS Code

Test that Copilot correctly reads and follows both instruction layers by checking

Generate minimal, targeted diffs

Practice using attached context (selections, symbols, files)

Governance and Scale-Out

Date: June 11, 2030



MCP Servers For Team Collaboration

Remote Deployment and Repo-scoped Config



Knowledge Sharing

MCP Servers shared across teams allow to share prompts, resources, tool implementations



Unified Setup

Every member of the team will have the same setup in seconds



Governance

ownership, audit, and access control in one place



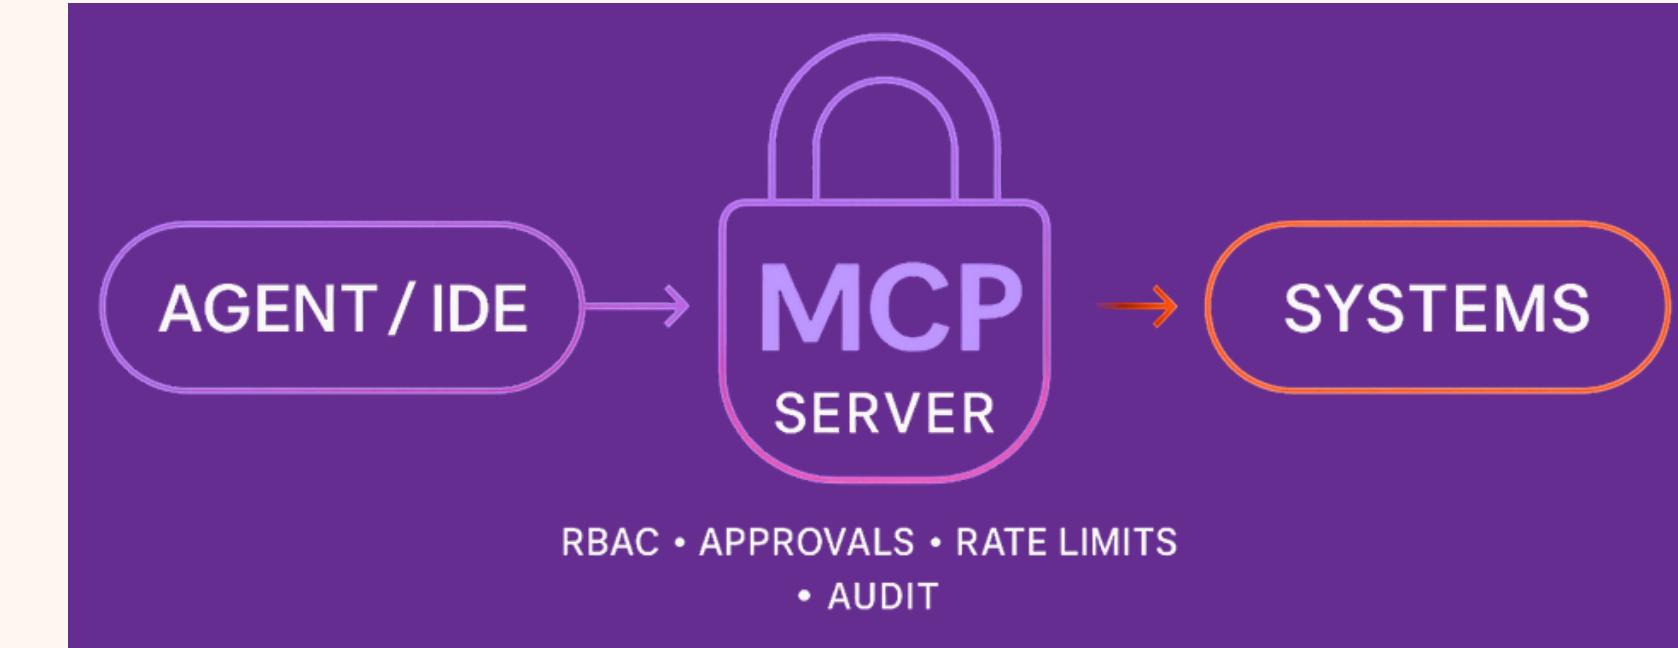
Faster Onboarding

Shared MCP servers give new teammates instant access to the same prompts, resources, and tools the rest of the org uses

Why Governance for MCP

Put controls where the power is - on the server.

-  **Client allowlists ≠ security:** enforce policy server-side
-  **Scale needs repeatability:** least-privilege access, standardized tool contracts, consistent rollouts.
-  **Auditability = trust:** log who/what/when; measure cost, latency, and errors.

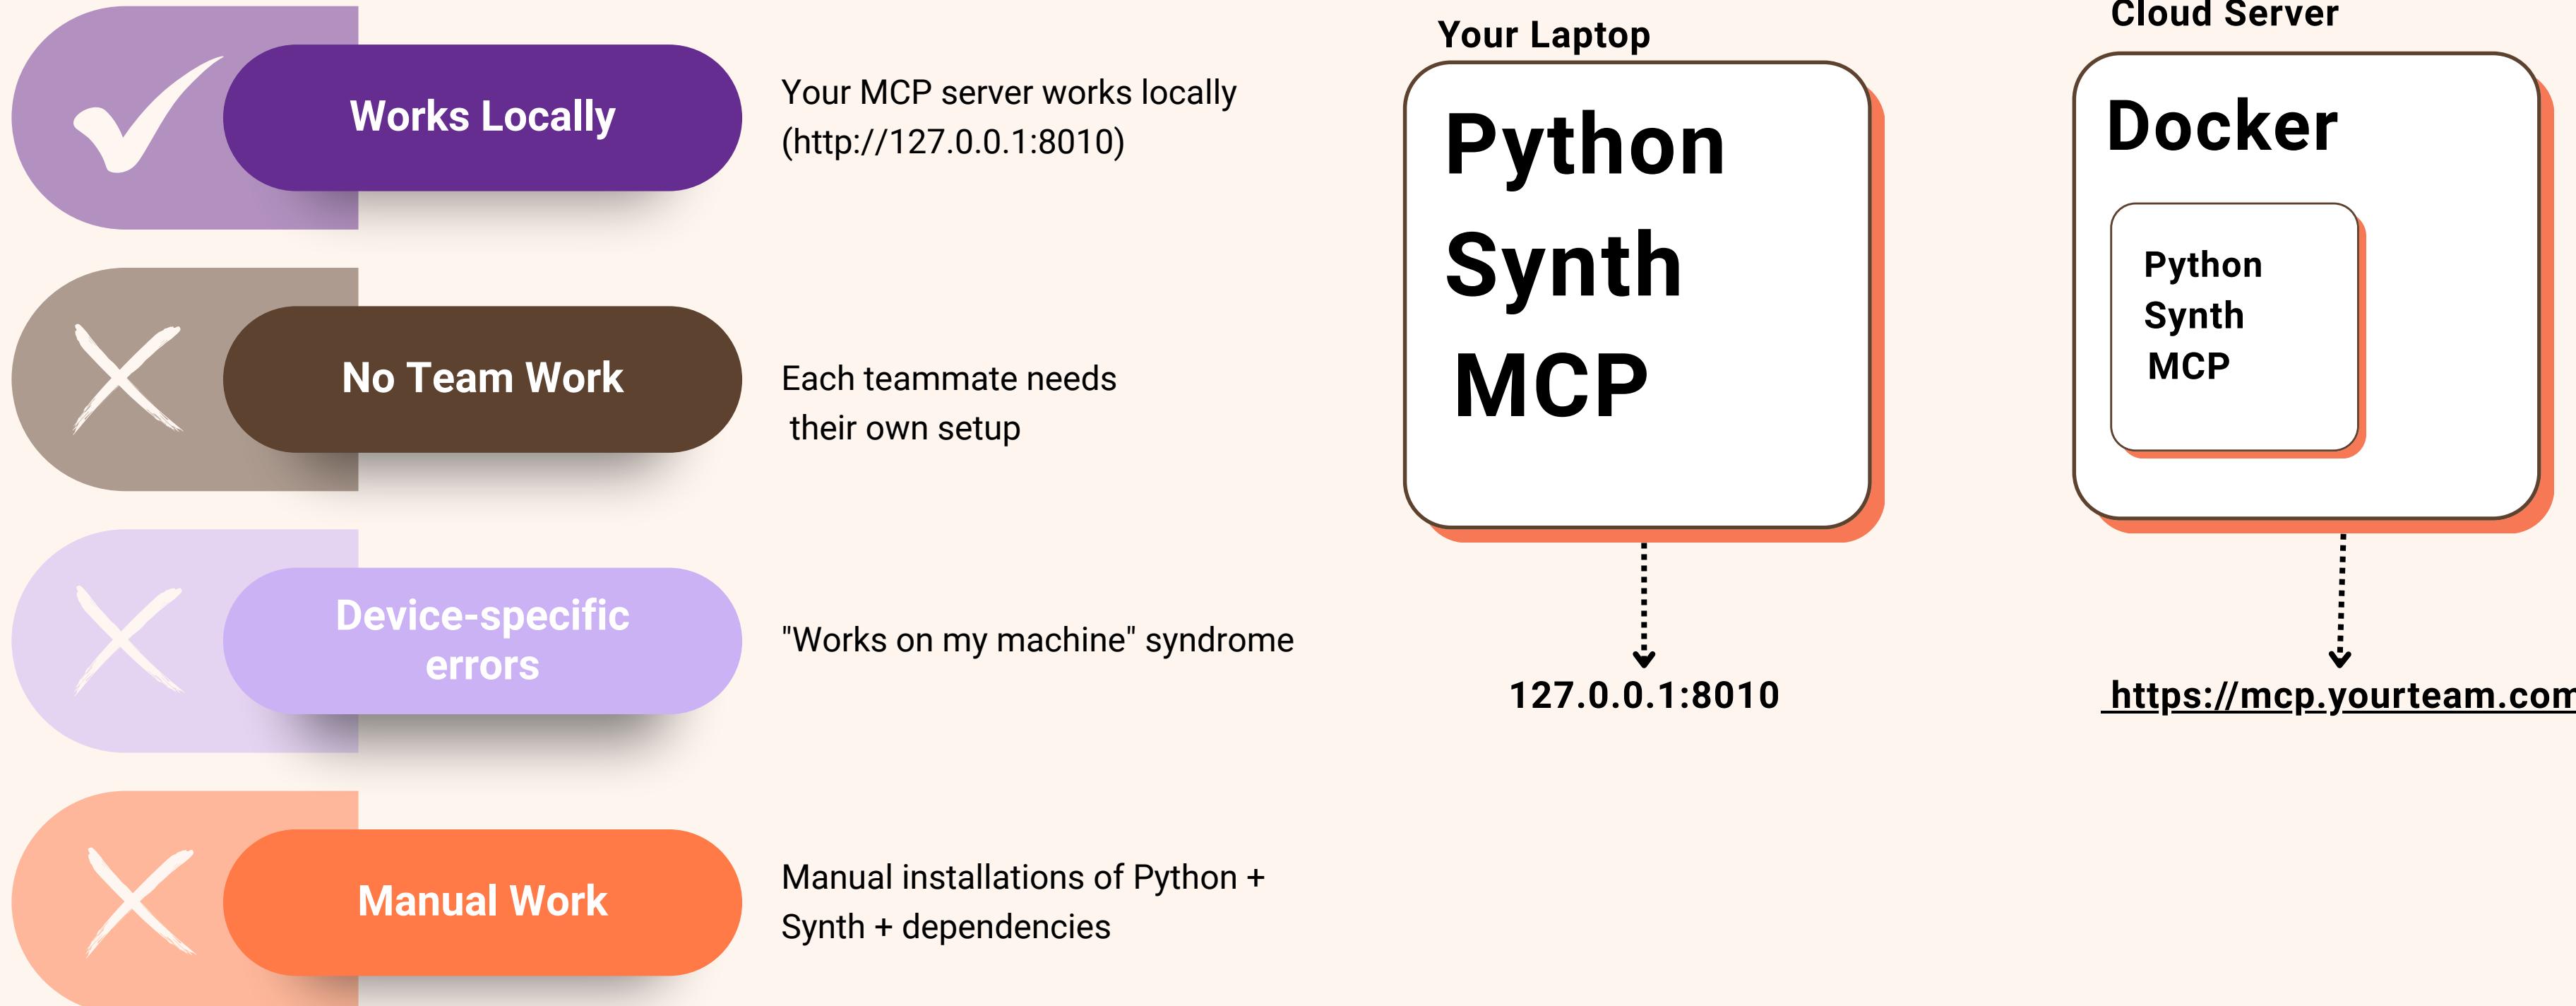


Empower teams without risking data, cost, or uptime.

We don't rely on the client's allowlist—governance lives at the MCP server with roles, approvals, caps, and audit so we can move fast safely.

From Local to Team-Wide:

Remote MCP Server Deployment



Minimal MCP Governance

Best Practices

```
{
  "servers": {
    >Start | More...
    "github-mcp": {
      "url": "https://api.githubcopilot.com/mcp/",
      "type": "http"
    },
    >Error | Restart
    "synth-local": {
      "type": "http",
      "url": "http://127.0.0.1:8010/mcp",
      "toolAllowList": ["synth_generate", "synth_inspect_model", "preview_tal
    }
  },
}
```

No authorisation for the server of synth-local
anyone can just call the server if they know the API

Unrestricted GitHub MCP
you might want to configure which tool permissions users have for your workspace github

No environment separation or secret handling
Hardcoded URLs as well as no separation of dev/stage/prod.

Client allowlist ≠ security
"toolAllowList" only hides tools in the UI but does not enforce on the server. You need server-side allow/deny, size/time caps, approvals, and rate limits..

Minimal MCP Governance

Best Practices

```
{
  "servers": {
    "github-mcp": {
      "type": "http",
      "url": "https://api.githubcopilot.com/mcp/",
      "toolAllowList": ["list_issues", "create_issue"]
    },
    "synth-remote": {
      "type": "http",
      "url": "${env:SYNTH_MCP_URL}/mcp",
      "headers": { "Authorization": "Bearer ${env:SYNTH_MCP_TOKEN}" },
      "toolAllowList": [
        "synth_generate",
        "synth_inspect_model",
        "preview_table_head"
      ]
    }
  }
}
```

No authorisation for the server of synth-local
anyone can just call the server if they know the API

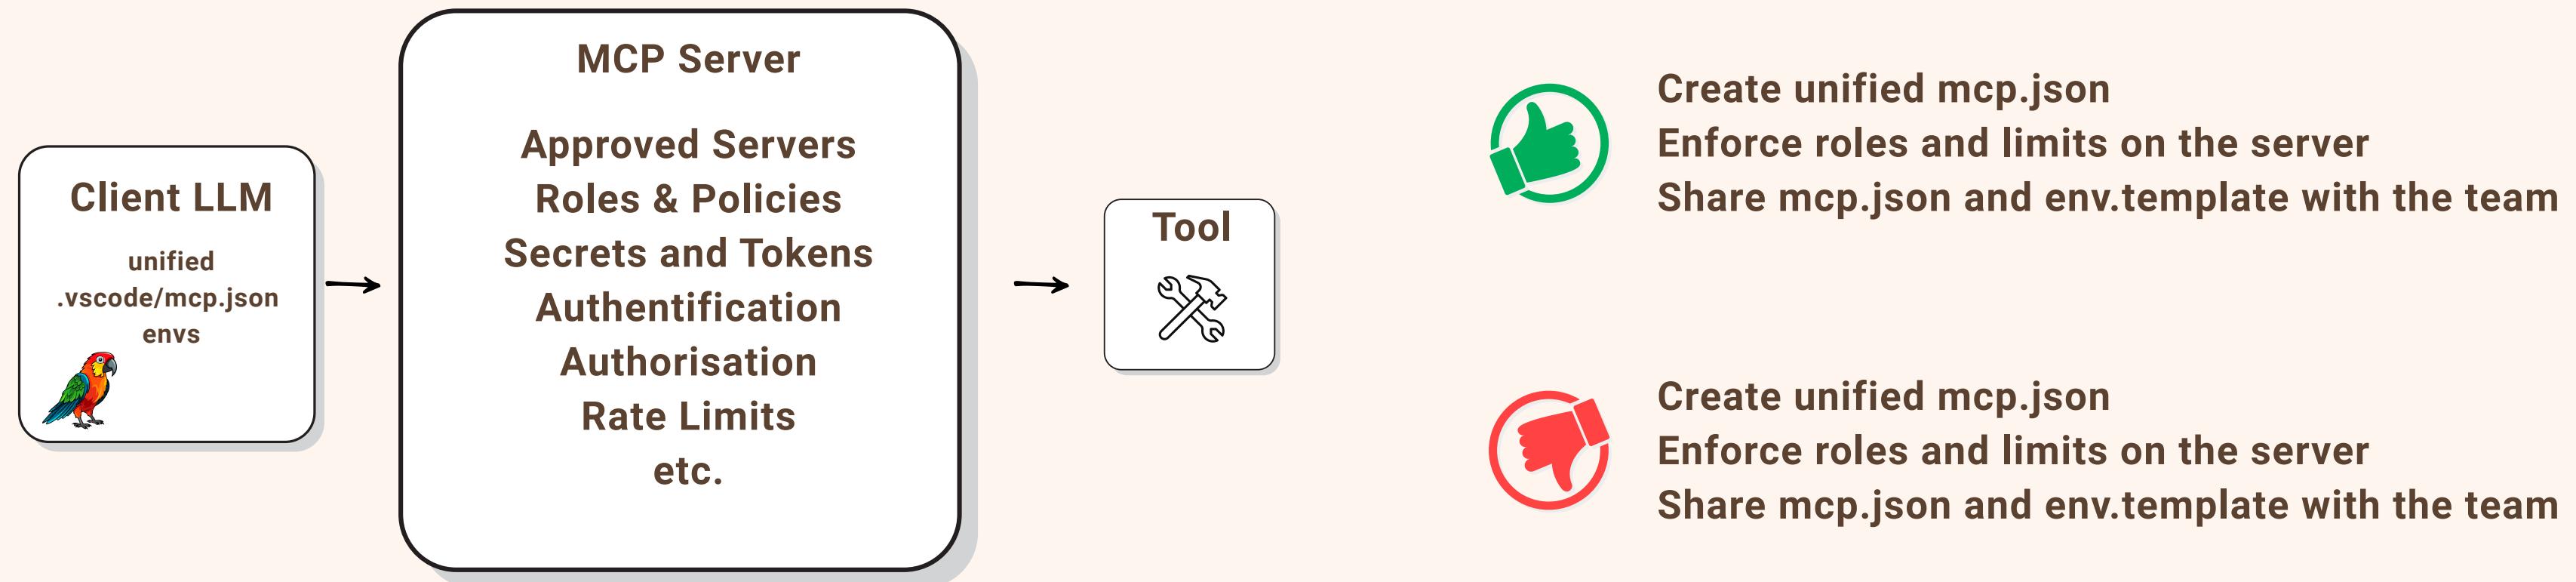
Unrestricted GitHub MCP
you might want to configure which tool permissions users have for your workspace github

No environment separation or secret handling
Hardcoded URLs as well as no separation of dev/stage/prod.

Client allowlist ≠ security
"toolAllowList" only hides tools in the UI but does not enforce on the server. You need server-side allow/deny, size/time caps, approvals, and rate limits..

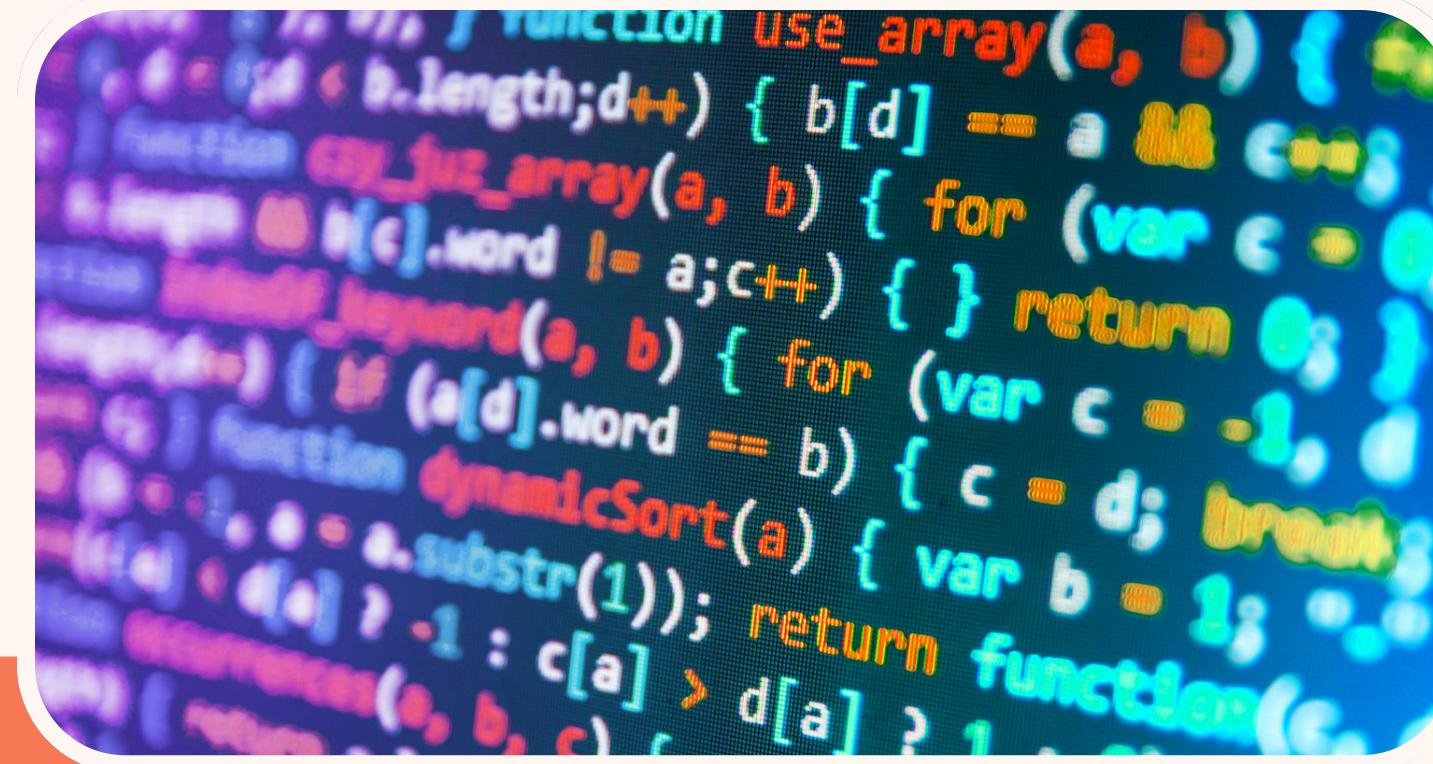
Team Collaboration

Best Practices



Conclusion

Coding Copilots, MCP servers etc. are still being developed and evolve significantly. Focus on understanding principals and not on UX.



Covered Topics:

- How to ground copilot in context
- How to create and use MCP servers
- How to generate unit tests
- How to govern MCP servers



Thank You For Your Attention



LinkedIn
[msukhareva](https://www.linkedin.com/in/msukhareva)



Email
info@mariasukhareva.com



Website
www.mariasukhareva.com

Presented by **Maria Sukhareva**

